

PThreads

A Note About Makefiles

- In OneFS/BSD be sure to build using `gmake`
- You may have to remove `-pedantic-errors` from `CFLAGS`

Creating/Starting a Thread (1 of 4)

```
#include <pthread.h>
typedef void *THREAD_PROC_t( void * );
typedef THREAD_PROC_t *THREAD_PROC_p_t;
int pthread_create(
    pthread_t          *thread,
    const pthread_attr_t *attr,
    THREAD_PROC_p_t     start_routine,
    void                *arg
);
```

thread

Pointer to a private thread ID.

attr

Pointer to an thread attributes block; may be NULL.

start_routine

Pointer to a function to control the thread.

arg

Argument passed to start_routine.

Creating/Starting a Thread (2 of 4)

Returns:

0 if successful, non-zero otherwise.

`pthread_create` creates and launches a thread by calling `start_routine` and passing `arg`.

```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <unistd.h>
#include <pthread.h>
```

```
static void *thread_proc( void * );
```

Creating/Starting a Thread (3 of 4)

```
int main( int argc, char **argv )
{
    pthread_t    thread;
    int stat =
        pthread_create( &thread, NULL, thread_proc, NULL );
    if ( stat != 0 )
    {
        const char *reason = strerror( stat );
        fprintf( stderr,
            "thread create failure, %d: \"%s\\",
            stat,
            reason
        );
        abort();
    }
    stat = pthread_join( thread, NULL );
    if ( stat != 0 )
    {
        . . .
    }
    return 0;
}
```

Creating/Starting a Thread (4 of 4)

```
static void *thread_proc( void *arg )
{
    for ( int inx = 0 ; inx < 2 ; ++inx )
    {
        puts( "thread" );
        sleep( 2 );
    }

    return NULL;
}
```

Waiting For Thread Completion

```
#include <pthread.h>
int pthread_join(pthread_t thread, void **retval);
```

thread

Thread ID.

retval

Pointer for returning thread termination value; may be NULL.

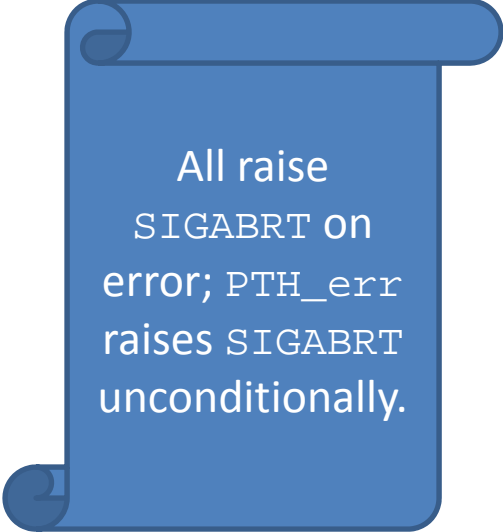
Returns:

0 if successful, non-zero otherwise.

Returns immediately if thread has already expired, otherwise returns after thread expiration.

The PTH Module

```
#include <pth.h>
void PTH_create(
    pthread_t          *thread,
    const pthread_attr_t *attr,
    PTH_PROC_p_t       start_routine,
    void               *arg
);
void PTH_join(
    pthread_t thread,
    void **retval
);
void PTH_err(
    int      status,
    const char *msg,
    const char *file,
    int line
);
```



All raise
SIGABRT on
error; PTH_err
raises SIGABRT
unconditionally.

See also:

`pthread_tryjoin_np`
`pthread_attr_init`
`pthread_attr_destroy`
`pthread_getattr_np`
`pthread_attr_*`

Exercises

1. Download the `pth` module from the class web site. Complete the functions `PTH_create` and `PTH_join`. Add the module to `libisi.a`.
2. Download `thr_test1.c` from the class web site. Complete the program so that `main` creates two threads based on `proc1` and `proc2`, and then joins the threads. Run the program.
3. For this exercise you will revise exercise 1 so that it runs both threads based on a single `start_routine`. Download `thr_test2.c` from the class web site. Complete `main` and `proc` according to the enclosed instructions. Run the program.
4. Download `thr_test3.c` from the class web site. Complete `main` according to the enclosed instructions. Run the program, redirecting `stdout` to a text file; `grep` the text file for lines containing “invalid.” What went wrong?

Race Conditions

THREAD1

```
while ( nextID_ < TEST_LEN )
{
    { int temp = nextID_;
      test_[nextID_] = nextID_;
    }
    { ++nextID_;
      usleep( uWaitTime_ );
    }
}
```

THREAD2

```
while ( nextID_ < TEST_LEN )
{
    int temp = nextID_;
    { test_[nextID_] = nextID_;
      ++nextID_;
      usleep( uWaitTime_ );
    }
}
```

1. nextID = 50
2. test_[50] = 50
3. temp = 50
4. nextID_ = 51
5. temp_[51] = 50, nextID_=52

Managing Race Conditions: Mutexes

```
#include <pthread.h>
typedef pthread_mutex_t ...
#define PTHREAD_MUTEX_INITIALIZER ...
int pthread_mutex_lock(pthread_mutex_t *mutex);
int pthread_mutex_trylock(pthread_mutex_t *mutex);
int pthread_mutex_unlock(pthread_mutex_t *mutex);
```

`pthread_mutex_lock`

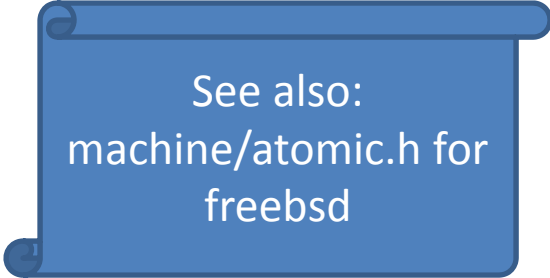
Locks a mutex. If the mutex is already locked, the thread is suspended until it is unlocked.

`pthread_mutex_trylock`

Trys to lock a mutex. If the mutex is already locked, returns a value of `EBUSY`.

`pthread_mutex_unlock`

Unlocks a mutex.

A blue callout box with rounded corners and a drop shadow, containing text.

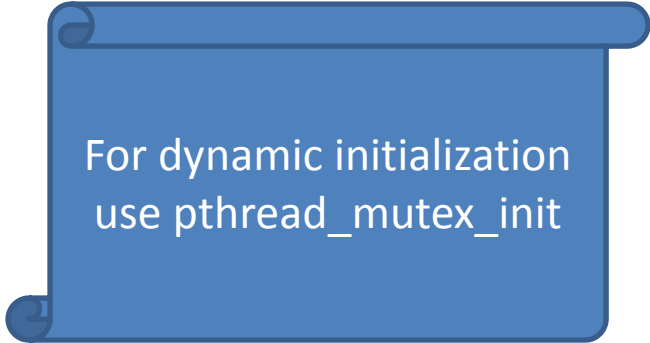
See also:
`machine/atomic.h` for
freebsd

See also:

`pthread_mutex_destroy`
`pthread_mutex_timedlock`
`pthread_mutexattr_init`
`pthread_mutexattr_destroy`
`pthread_mutexattr_*`

Using a Mutex

```
static pthread_mutex_t  mutex_          =  
    PTHREAD_MUTEX_INITIALIZER;  
static int              store_inx_      = 0;  
static int              storage_[100];  
  
// May be called from multiple threads  
static void store_val( int val )  
{  
    pthread_mutex_lock( &mutex_ );  
    if ( store_inx_ < ISI_CARD( storage_ ) )  
        storage_[store_inx_++] = val;  
    pthread_mutex_unlock( &mutex_ );  
}
```

A blue callout box with rounded corners and a drop shadow, containing text about dynamic initialization.

For dynamic initialization
use `pthread_mutex_init`

Exercises

5. Revise `thr_test3.c` so that it uses a mutex to fix the race condition.

Threading 1-Step Processes For Multiple Clients

```
do
{
    client = nextClient()
    if ( client != NULL )
    {
        dispatch client via new thread
        store thread
    } while ( client != NULL )
sequentially join all threads
```


Threading N-Step Processes

```
wait for new component  
post component to step 1 queue
```

Initiator thread

```
wait for next component  
process component  
post component to step 2 queue
```

Step 1 thread

```
wait for next component  
process component  
post component to step 3 queue
```

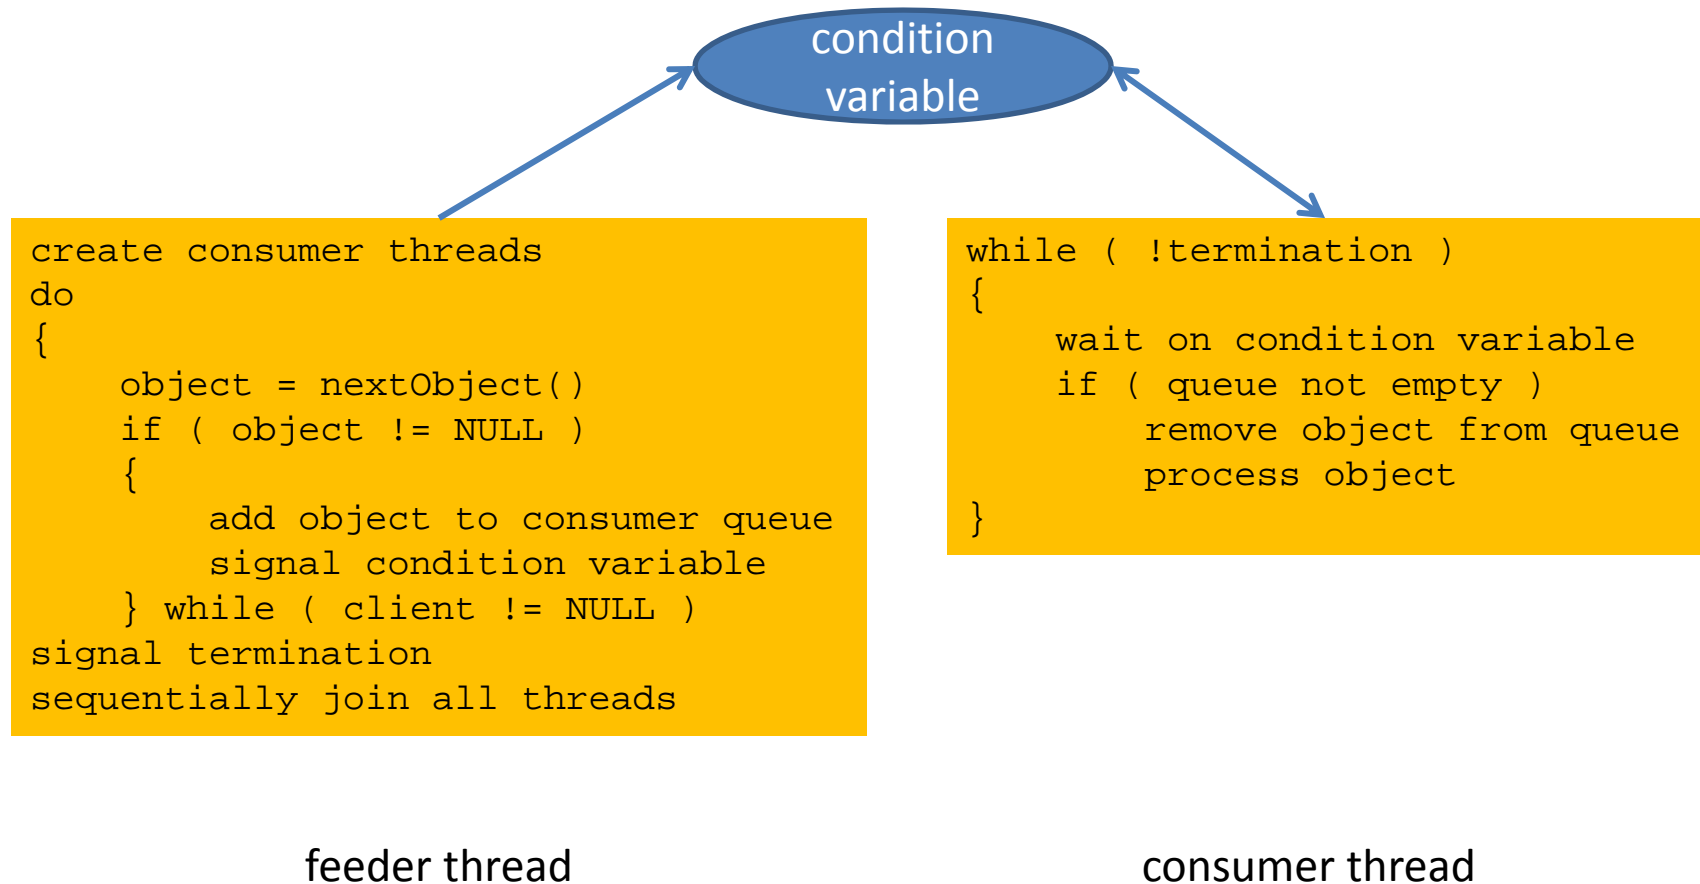
Step 2 thread

. . .

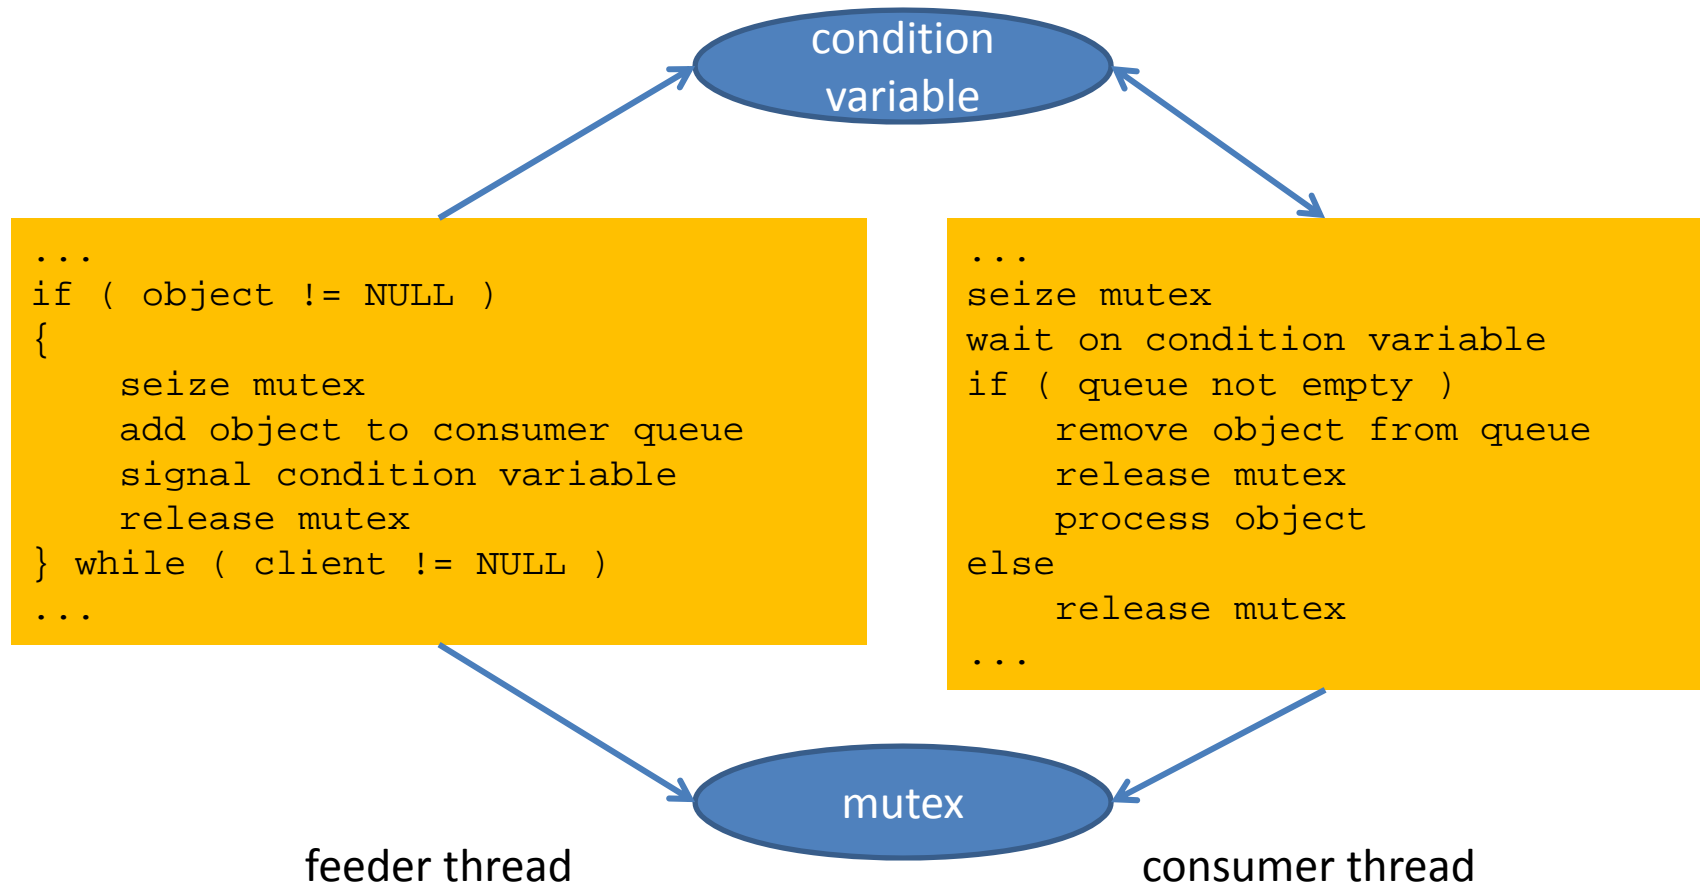
```
wait for next component  
dispose component
```

Finalizer thread

Waiting: Condition Variables



Synchronizing Condition Variable Usage



pthread_cond_wait

```
int pthread_cond_wait(  
    pthread_cond_t *cond,  
    pthread_mutex_t *mutex  
);
```

- Call wait function with `mutex` locked
- If `cond` has been signaled wait returns immediately...
- ... else thread is suspended and `mutex` is unlocked
- Thread is always wakened with `mutex` locked

pthread_cond_timedwait

```
int pthread_cond_timedwait(  
    pthread_cond_t      *cond,  
    pthread_mutex_t      *mutex,  
    const struct timespec *abstime  
);
```

- Like `pthread_cond_wait` except may return `ETIMEDOUT`
- `abstime` is the time to wake the thread, *not* an interval

struct timespec

```
#include <time.h>
struct timespec
{
    time_t tv_sec;    /* seconds */
    long    tv_nsec;  /* nanoseconds */
};
```

tv_nsec must be in the range 0 to 999,999,999

pthread_cond_signal

```
int pthread_cond_broadcast(pthread_cond_t *cond);  
int pthread_cond_signal(pthread_cond_t *cond);
```

- `pthread_cond_signal`: If more than one thread is waiting on `cond`, only one is wakened.
- `pthread_cond_broadcast`: All threads waiting on `cond` are wakened.

Condition Variable Example (1 of 8)

```
typedef struct control_s
{
    pthread_mutex_t      mutex;
    pthread_cond_t       cond;
    int                  ident;
} CONTROL_t, *CONTROL_p_t;

static void *thrProc( void * );

static int              inProgress_ = 1;
static CONTROL_t        controls[2];
```


Condition Variable Example (2 of 8)

```
int main( int argc, char **argv )
{
    pthread_t    thr1;
    pthread_t    thr2;

    PTH_mutex_init( &controls_[0].mutex, NULL );
    PTH_cond_init( &controls_[0].cond, NULL );
    controls_[0].ident = 1;
    PTH_create( &thr1, NULL, thrProc, &controls_[0] );

    PTH_mutex_init( &controls_[1].mutex, NULL );
    PTH_cond_init( &controls_[1].cond, NULL );
    PTH_create(&thr2, NULL, thrProc, &controls_[1]);
    controls_[1].ident = 2;
    puts( "start" );
    . . .
}
```

Condition Variable Example (3 of 8)

```
int main( int argc, char **argv )
{
    . . .
    for ( int inx = 0 ; inx < 5 ; ++inx )
    {
        sleep( 2 );
        pthread_mutex_lock( &controls_[0].mutex );
        pthread_cond_signal( &controls_[0].cond );
        pthread_mutex_unlock( &controls_[0].mutex );
        if ( inx %2 == 0 )
        {
            PTH_mutex_lock( &controls_[1].mutex );
            PTH_cond_signal( &controls_[1].cond );
            PTH_mutex_unlock( &controls_[1].mutex );
        }
    }
    . . .
}
```

Condition Variable Example (4 of 8)

```
    . . .  
    inProgress_ = 0;  
    PTH_join( thr1, NULL );  
    PTH_join( thr2, NULL );  
    puts( "end" );  
  
    return 0;  
  
}
```

Condition Variable Example (5 of 8)

```
static void *thrProc( void *arg )  
{  
    CONTROL_p_t control = arg;  
    . . .
```

Condition Variable Example (6 of 8)

```
. . .  
while ( inProgress_ )  
{  
    struct timespec waitTime = { time( NULL ) + 2 , 0 };  
  
    PTH_mutex_lock( &control->mutex );  
    errno = 0;  
    PTH_cond_timedwait( &control->cond,  
                        &control->mutex,  
                        &waitTime  
                        );  
  
    const char *msg =  
        errno == ETIMEDOUT ? "timeout" : "signal";  
    printf("Thread %d: %s\n", control->ident, msg );  
    PTH_mutex_unlock( &control->mutex );  
}  
. . .
```



... about
two
seconds
from now

Condition Variable Example (7 of 8)

```
    . . .  
    printf( "Thread %d exiting\n", control->ident );  
  
    return NULL;  
}
```

Condition Variable Example (8 of 8)

```
$ ./thr_test5
start
Thread 1: timeout
Thread 2: timeout
Thread 1: signal
Thread 2: signal
Thread 1: timeout
Thread 2: timeout
Thread 1: signal
Thread 2: timeout
. . .
Thread 1: timeout
Thread 2: timeout
Thread 1: signal
Thread 2: timeout
Thread 1: timeout
Thread 2: timeout
Thread 1: signal
Thread 1 exiting
Thread 2: signal
Thread 2 exiting
end
```

See also:

- `pthread_cond_destroy`
- `pthread_condattr_init`
- `pthread_condattr_destroy`
- `pthread_condattr_*`

Exercises

6. Download `cond_exercise.c` from the class website. This program uses three threads to (1) obtain an item to process (`thrGet/procGet`), (2) process the item (`thrOne/procOne`), and (3) finalize the processing (`thrFin/procFin`). `thrGet` posts the item to `anchorOne`, synchronizing via `mutexOne` and `condOne`. When ready, `thrOne` posts the item to `anchorFin`, synchronizing via `mutexFin` and `condFin`. Add another step to this process. Create an anchor, mutex and condition variable for `thrTwo/procTwo`. Change `procOne` so that it posts to the queue for `thrTwo`, and write `procTwo` so that it processes the item (via `work()`) and posts it to `anchorFin`. Create a new thread in `main` to drive `thrTwo`.
7. Modify `cond_exercise.c`. Add an object (`struct`) that can contain an anchor, mutex and condition variable. Use this object to encapsulate the anchor, mutex and condition variable for each of the four threads. You will probably have to initialize the mutex and condition variable via `PTH_mutex_init` and `PTH_cond_init`.
8. Design an object that can encapsulate an anchor, mutex, condition variable, pointer to thread function, thread ID, process ID (type `int`) and a `struct timespec` object.

Synchronizing via Signals

Threads can communicate via signals.

- The signals to communicate with must be *blocked*; use `pthread_sigmask`.
- Signal the thread using `pthread_kill`.
- The thread waits for the signal using `sigwait` or `sigwaitinfo`.

pthread_sigmask (1)

```
#include <pthread.h>
#include <signal.h>
int pthread_sigmask(
    int          how,
    const sigset_t *set,
    sigset_t      *oldset
);
```

how

use SIG_BLOCK to block, SIG_UNBLOCK to unblock;
may be NULL.

set

signals affected by this call.

oldset

previous signal mask; may be NULL.

pthread_sigmask (2)

```
sigset_t    mask;  
sigemptyset( &mask );  
sigaddset( &mask, SIGHUP );  
sigaddset( &mask, SIGUSR1 );  
PTH_sigmask( SIG_BLOCK, &mask, NULL );
```

pthread_kill

```
#include <pthread.h>
#include <signal.h>
int pthread_kill(
    pthread_t thread,
    int      sig
);
```

thread
the thread to signal.

sig
the signal to send.

sigwait (1)

```
#include <signal.h>
int sigwait(
    sigset_t set,
    int      *sig
);
```

set

the signals to wait on.

sig

the signal that ended the wait.

return

0 for success, error number otherwise.

sigwait (2)

```
sigset_t    sigset;
sigemptyset( &sigset );
sigaddset( &sigset, SIGHUP );
sigaddset( &sigset, SIGUSR1 );

int stat     = sigwait( &sigset, &sig );
if ( stat != 0 )
{
    const char *msg     = strerror( stat );
    fprintf( stderr, "sigwait: \"%s\"\n", msg );
    abort();
}
```

`sigwait` (3)

- If a signal is pending returns immediately
- If a signal has been posted more than once when `sigwait` is called, behavior is system-dependent: one signal may be cleared or all signals may be cleared

See also:

`sigwaitinfo`
`sigtimedwait`

sigwait Example (1)

```
static void init( void )
{
    sigset_t    mask;
    sigemptyset( &mask );
    sigaddset( &mask, SIGHUP );
    sigaddset( &mask, SIGUSR1 );
    PTH_sigmask( SIG_BLOCK, &mask, NULL );
}
```

sigwait Example (2)

```
static void *thrProc( void *arg )
{
    sleep( 2 );
    sigset_t    sigset;
    int         sig      = SIGUSR1;
    int         count    = 0;

    sigemptyset( &sigset );
    sigaddset( &sigset, SIGHUP );
    sigaddset( &sigset, SIGUSR1 );
    . . .
```

sigwait Example (3)

```
    . . .  
do  
{  
    int stat      = sigwait( &sigset, &sig );  
    if ( stat != 0 )  
    {  
        const char *msg      = strerror( stat );  
        fprintf( stderr, "%s\n", msg );  
        abort();  
    }  
    if ( sig == SIGUSR1 )  
        printf( "SIGUSR1: %d\n", ++count );  
    if ( sig == SIGHUP )  
        printf( "SIGHUP\n" );  
} while ( sig == SIGUSR1 );  
  
return NULL;  
}
```

sigwait Example (4)

```
int main( int argc, char **argv )
{
    init();
    pthread_t  thr;
    printf( "start\n" );
    PTH_create( &thr, NULL, thrProc, NULL );
    PTH_kill( thr, SIGUSR1 );
    PTH_kill( thr, SIGUSR1 );
    PTH_kill( thr, SIGUSR1 );
    sleep( 3 );
    PTH_kill( thr, SIGUSR1 );
    sleep( 3 );
    PTH_kill( thr, SIGUSR1 );
    PTH_kill( thr, SIGUSR1 );
    sleep( 3 );
    PTH_kill( thr, SIGHUP );
    PTH_join( thr, NULL );
    printf( "end\n" );

    return 0;
}
```

sigwait Example (5)

```
$ ./thr_sigwait  
start  
SIGUSR1: 1  
SIGUSR1: 2  
SIGUSR1: 3  
SIGHUP  
end
```

Exercises

9. Download `sigwait_exercise.c`. Complete the code as instructed in the file.

Read/Write Locks

For a resource, a read/write lock:

- May be seized for read or write access, and must subsequently be released
- Maintains a read count and a write count
- When write count is 0, unlimited read access is allowed
- When write count is 1, all access is blocked until the read count is 0, then the write access is allowed

pthread_rwlock_...

```
pthread_rwlock_t rwlock =
    PTHREAD_RWLOCK_INITIALIZER;
int pthread_rwlock_init(
    pthread_rwlock_t *rwlock,
    pthread_rwlockattr_t *attr
);
int pthread_rwlock_rdlock(
    pthread_rwlock_t *rwlock
);
int pthread_rwlock_tryrdlock(
    pthread_rwlock_t *rwlock
);
int pthread_rwlock_wrlock(
    pthread_rwlock_t *rwlock
);
int pthread_rwlock_unlock(
    pthread_rwlock_t *rwlock
);
```


PTH_rwlock_...

```
...  
void PTH_rwlock_tryrdlock(  
    pthread_rwlock_t *rwlock  
);
```

If the operation fails with a status of `EPERM`, `errno` is set to `EPERM`.

Read/Write Lock Example (1)

```
typedef struct control_s
{
    int          pauseBefore;
    int          pauseAfter;
    int          type;
    int          ident;
    pthread_t     thrID;
} CONTROL_t, *CONTROL_p_t;
static pthread_rwlock_t rwlock =
    PTHREAD_RWLOCK_INITIALIZER;
. . .
```

Read/Write Lock Example (2)

```
CONTROL_t  threads[]  =
{
    { 0000, 4000, TYPE_READ,  0 },
    { 1000, 4000, TYPE_READ,  0 },
    { 2000, 4000, TYPE_READ,  0 },
    { 3000, 5000, TYPE_WRITE, 0 },
    { 4000, 2000, TYPE_READ,  0 },
    { 5000, 2000, TYPE_READ,  0 },
};

for ( int inx = 0 ;
      inx < ISI_CARD( threads ) ; ++inx )
{
    threads[inx].ident = inx;
    PTH_create( &threads[inx].thrID, NULL,
                driver, &threads[inx] );
}
```

Read/Write Lock Example (3)

```
static void *driver( void *arg )
{
    CONTROL_p_t contr    = arg;
    pause( contr->pauseBefore );
    switch ( contr->type )
    {
        case TYPE_READ:
            printf( "... rdlock ..." );
            PTH_rwlock_rdlock( &rwlock );
            break;
        case TYPE_WRITE:
            printf( "... write lock ..." );
            PTH_rwlock_wrlock( &rwlock );
            break;
    }
    printf( ".. lock acquired ..." );
    pause( contr->pauseAfter );
    printf( "... unlocking\n" );
    PTH_rwlock_unlock( &rwlock );
}
```

Read/Write Lock Example (4)

```
$ ./rwlock_test2
thread 0 attempting read_lock
thread 0 lock acquired
    thread 1 attempting read_lock
    thread 1 lock acquired
        thread 2 attempting read_lock
        thread 2 lock acquired
            thread 3 attempting write_lock
            thread 4 attempting read_lock
thread 0 unlocking
    thread 1 unlocking
        thread 5 attempting read_lock
thread 2 unlocking
    thread 3 lock acquired
    thread 3 unlocking
        thread 4 lock acquired
            thread 5 lock acquired
            thread 5 unlocking
            thread 4 unlocking
```

Preventing Writer Starvation (1)

When read-locks and write-locks are blocked, priority is given to write-locks when unblocked.

```
CONTROL_t    threads[ ]    =
{
    { 0000, 4000, TYPE_WRITE, 0 },
    { 1000, 4000, TYPE_READ,  0 },
    { 2000, 4000, TYPE_READ,  0 },
    { 3000, 5000, TYPE_WRITE, 0 }
};
```

Preventing Writer Starvation (2)

```
$ ./rwlock_test3
thread 0 attempting write_lock
thread 0 lock acquired
    thread 1 attempting read_lock
        thread 2 attempting read_lock
            thread 3 attempting write_lock
thread 0 unlocking
    thread 3 lock acquired
        thread 3 unlocking
            thread 2 lock acquired
thread 1 lock acquired
    thread 2 unlocking
thread 1 unlocking
```

nanosleep

```
int nanosleep(  
    struct timespec *req,  
    struct timespec *rem  
);  
struct timespec {  
    time_t tv_sec;           /* seconds */  
    long   tv_nsec;         /* nanoseconds */  
};
```

Sleeps for the number of seconds/nanoseconds indicated by `req`. If the operation is interrupted, `errno` is set to `EINTR`. and the remaining (unelapsed) time is returned in `rem` (if it is non-NULL).

See also:

`pthread_rwlock_destroy`
`pthread_rwlockattr_init`
`pthread_rwlockattr_destroy`
`pthread_rwlockattr_*`

Exercises

10. Write a program that launches eleven threads; ten of them should wait one quarter of a second; lock a read-write mutex for reading; wait another quarter of a second; release the mutex; and repeat. The eleventh should wait two seconds; lock the mutex for writing; wait another quarter of a second; release the mutex; and repeat. Each thread should keep a count of the number of times a lock succeeded. Use a global flag to notify your threads when to quit. When all of the threads have been terminated, your main thread should print out the counts for each thread. The program should run for 15 seconds.

Suggestions: a) write two start-routines, one for the readers and one for the writer. b) use a `struct` to store the thread counts and thread IDs; use a pointer to the `struct` as the start-routine argument.

Spin Locks

- A spin lock is a compute-intensive type of lock
- Spin locks are intended for multiprocessor environments
- Spin locks are rarely used in application-level code
- Spin locks should be used with great care

```
static volatile bool available_;  
while ( !available_ )  
    ;  
available = false;  
execute();
```

See also:

```
void PTH_spin_init(  
    pthread_spinlock_t *lock,  
    int pshared  
);  
void PTH_spin_destroy(  
    pthread_spinlock_t *destroy  
);  
void PTH_spin_lock(  
    pthread_spinlock_t *lock  
);  
void PTH_spin_unlock(  
    pthread_spinlock_t *lock  
);
```

Critical Sections

- A critical section is a block of code that must be executed by at most one thread at a time
- Critical sections can be implemented using mutexes
- Code in critical sections should *not* perform blocking operations

```
static pthread_mutex_t  critSecMutex_  =  
    PTHREAD_MUTEX_INITIALIZER;
```

```
void thread_funk( void )  
{  
    PTH_mutex_lock( &critSecMutex_ );  
    . . .  
    PTH_mutex_unlock( &critSecMutex_ );  
}
```

Thread Management – Pay As You Go

- Create a new thread every time you need one
- Thread must be detached
- Once a thread is detached it can't be joined

```
while ( !shutDown_ )
{
    pthread_t thread;
    DATA_      *data = nextClient();
    PTH_create( &thread, NULL, driver, data );
    PTH_detach( thread );
}
```

pthread_detach

```
int pthread_detach(  
    pthread_t thread_id  
);
```

Normally thread resources for a thread are not released until the thread is joined. Use detach when you don't expect to join with a thread.

- You cannot join a detached thread.
- You cannot communicate with the thread via its ID.
- All thread resources are freed when the thread exits.

Problems With This Approach

- Since threads can't be joined you need a way to know when all threads have exited.
- You may need an infinite supply of threads

```
while ( !shutDown_ )
{
    pthread_t thread;
    DATA_ *data = nextClient();
    PTH_create( &thread, NULL, driver, data );
    PTH_detach( thread );
}
```


Tracking Detached Thread Progress

- Since threads can't be joined you need a way to know when all threads have exited.
- You may need an infinite supply of threads

```
static int num_threads_    = 0;
. . .
while ( (data = nextClient()) == NULL )
{
    nextThread( data );
}

while ( num_threads_ > 0 )
    pauseMillis( 125 );
```

Governing Thread Creation (1 of 3)

```
static void nextThread( int *arg )
{
    PTH_mutex_lock( &mutex_ );
    ISI_ASSERT( num_threads_ <= MAX_THREADS );
    ISI_ASSERT( num_threads_ >= 0 );

    bool    done    = false;
    while ( !done )
    {
        if ( num_threads_ == MAX_THREADS )
        {
            printf( "waiting...\n" );
            PTH_cond_wait( &cond_, &mutex_ );
        }
        . . .
    }
```

Governing Thread Creation (2 of 3)

```
. . .  
    if ( num_threads_ < MAX_THREADS )  
    {  
        pthread_t  thread;  
        PTH_create(&thread, NULL, driver, arg);  
        PTH_detach( thread );  
        ++num_threads_;  
        done = true;  
    }  
}  
PTH_mutex_unlock( &mutex_ );  
}
```

Governing Thread Creation (3 of 3)

```
static void *driver( void *arg )
{
    int *ident  = arg;
    printf( "thread %2d starting\n", *ident );
    pauseMillis( 500 );

    PTH_mutex_lock( &mutex_ );
    --num_threads_;
    PTH_cond_signal( &cond_ );
    PTH_mutex_unlock( &mutex_ );

    printf( "thread %2d terminating\n", *ident );

    return NULL;
}
```

Problems With This Implementation

- The main thread may be blocked for a long time.

Solution:

- Create a separate thread to dispatch worker threads.
- The main (or some other) thread simply enqueues clients on a dispatch queue.

```
while ( !shutDown_ )
{
    pthread_t thread;
    DATA_ *data = nextClient();
    PTH_create( &thread, NULL, driver, data );
    PTH_detach( thread );
}
```

Problems With This Implementation

- The main thread may be blocked for a long time.

Solution:

- Create a separate thread to dispatch worker threads.
- The main (or some other) thread simply enqueues clients on a dispatch queue.

```
while ( !shutDown_ )
{
    pthread_t thread;
    DATA_ *data = nextClient();
    PTH_create( &thread, NULL, driver, data );
    PTH_detach( thread );
}
```

Exercises

11. Download `pay_as_you_go.c` from the class website. Modify it so that it creates a dispatcher thread which takes client data from a queue and uses it to create and detach a new thread.

Suggestions:

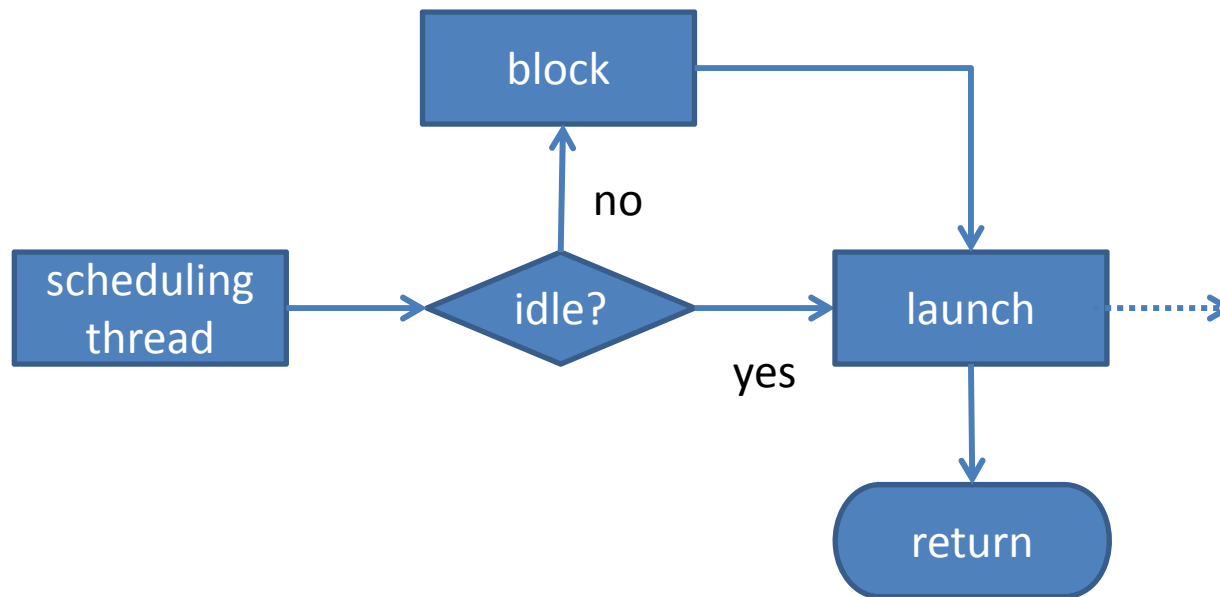
- a) Create an object that contains an anchor, mutex, condition variable and shutdown flag. Pass an object of this kind as the arg of the dispatcher thread.
- b) Create a subclass of `ENQ_ITEM_t` that contains a field for client data.
- c) After the while loop in `main` terminates, set the shutdown flag and join with the dispatcher thread.
- d) In the dispatcher thread create a loop that runs as long as the shutdown flag isn't set or the queue isn't empty.
- e) After the loop in the dispatcher thread terminates wait for all threads to complete then exit.

Thread Management – Thread Pools

- Designate a fixed number of persistent threads to manage all tasks
- Eliminates overhead associated with thread “churn”
- Two basic models: *task-blocked* vs. *thread-blocked*


Task-Blocked Thread Pools

- Idle threads reside on a stack
- A new task is assigned to thread at top of the stack
- If the stack is empty, the calling thread is blocked
- Advantage: efficient
- Disadvantage: calling thread may be blocked



Thread-Blocked Thread Pools

- Idle threads block on a queue
- New task is queued, queue is signaled
- Pool manager immediately returns
- Advantage: calling thread is never blocked
- Disadvantage: less efficient



Thread Pool API: `tpoolp.h` (1 of 2)

```
typedef struct tpool__control_s  
    *TPOOL__CONTROL_p_t;
```

```
typedef struct tpool__task_s  
{  
    ENQ_ITEM_t          item;  
    TPOOL_THREAD_PROC_p_t proc;  
    void               *arg;  
} TPOOL__TASK_t, *TPOOL__TASK_p_t;
```



forward
reference

Thread Pool API: tpoolp.h (2 of 2)

```
typedef struct tpool__thread_block_s  
{
```

```
    void          *arg;  
    TPOOL_THREAD_PROC_p_t proc;  
    pthread_t      thrID;  
    TPOOL__CONTROL_p_t control;
```

```
} TPOOL__THREAD_BLOCK_t, *TPOOL__THREAD_BLOCK_p_t;
```

declared in
tpool.h

```
typedef struct tpool__control_s  
{
```

```
    int          numThreads;  
    ENQ_ANCHOR_p_t tasks;  
    TPOOL__THREAD_BLOCK_p_t threads;  
    pthread_mutex_t flagMutex;  
    pthread_mutex_t taskMutex;  
    pthread_cond_t taskCond;  
    bool          quiesce;
```

```
} TPOOL__CONTROL_t;
```

Thread Pool API: tpool.h

```
#define TPOOL_NULL_ID    (NULL)
```

```
typedef struct tpool__control_s *TPOOL_ID_t;
```

```
typedef void *TPOOL_THREAD_PROC_t( void * );
```

```
typedef TPOOL_THREAD_PROC_t *TPOOL_THREAD_PROC_p_t;
```

Thread Pool API: Methods (1 of 2)

```
TPOOL_ID_t TPOOL_create(int numThreads );
```

Create a thread pool of size numThreads.

```
TPOOL_ID_t TPOOL_destroy(TPOOL_ID_t pool );
```

Destroy a thread pool.

```
void TPOOL_add_task(  
    TPOOL_ID_t          pool,  
    TPOOL_THREAD_PROC_p_t proc,  
    void                *arg  
);
```

Add a task to the task queue.

Thread Pool API: Methods (2 of 2)

```
void TPOOL_quiesce(TPOOL_ID_t pool );
```

Begin gradual shutdown; threads exit when all tasks are complete.

```
bool TPOOL_has_task(TPOOL_ID_t pool );
```

Returns true if the task queue is empty.

```
void TPOOL_join( TPOOL_ID_t pool );
```

Returns after every thread in the pool has been joined.

See Also

`threadpool.tgz`

On the class web site.

Author:

Tomer Heber

<http://sourceforge.net/projects/cthreadpool/>

`gdb:`

`info thread command`

`break command, thread parameter:`

`(gdb) b 37 thread 4`