



CSCI-GA.3033-012
**Multicore Processors:
Architecture & Programming**

Lecture 2: Concurrency and Parallelism

Mohamed Zahran (aka Z)


mzahran@cs.nyu.edu


<http://www.mzahran.com>




Same Meaning?

- **Concurrency**: At least two tasks are making progress at the same time frame.
 - Not necessarily at the same time
 - Include techniques like time-slicing
 - Can be implemented on a single processing unit
 - Concept more general than parallelism
- **Parallelism**: At least two tasks execute literally at the same time.
 - Requires hardware with multiple processing units





Concurrency without parallelism



Concurrency with parallelism

Performance tuning technique number 106: Concurrency vs. Parallelism

Copyright © Fasterj.com Limited

Questions!

If we have as much hardware as we want,
do we get as much parallelism as we wish?

If we have 2 cores, do we get 2x speedup?


Amdahl's Law




Gene M. Amdahl

- How much of a speedup one could get for a given parallelized task?


If F is the fraction of a calculation that is sequential then the maximum speed-up that can be achieved by using P processors is $1/(F+(1-F)/P)$




1CPU



2CPUs



3CPUs



4CPUs

What Was Amdahl Trying to Say?

- Don't invest blindly on large number of processors.
- Having faster core (or processor at his time) makes more sense than having many cores.


Was he right?

- At his days (the law appeared 1967) many programs have long sequential parts.
- This is not necessarily the case nowadays.
- It is not very easy to find F (sequential portion)

So ...

- Decreasing the serialized portion is of greater importance than adding more cores
- Only when a program is mostly parallelized, does adding more processors help more than parallelizing the remaining rest
- **Gustafson's law:** computations involving arbitrarily large data sets can be efficiently parallelized
- Both Amdahl and Gustafson do not take into account:
 - The overhead of synchronization, communication, OS, etc.
 - Load may not be balanced among cores
- So you have to use these laws as guideline and theoretical bounds only.

DAG Model for Multithreading




Work: total amount of time spent on all instructions

T_p = The fastest possible execution time on P processors

Work Law:

$$T_p \geq T_1/P$$

DAG Model for Multithreading



Span: The longest path of dependence in the DAG
 $= T_{\infty}$

Span Law:

$$T_p \geq T_{\infty}$$


Can We Define Parallelism Now?

How about?

$$T_1/T_\infty$$

Ratio of work to span

Can We Define Parallelism Now?



Work: $T_1 = 50$

Span: $T_\infty = 8$

Parallelism: $T_1/T_\infty = 6.25$


At What Level Can We Reason
About Parallelism
(algorithm, high-level language,
assembly)?

Is Thread The Only Parallelism Granularity?

- Instruction level parallelism (ILP)
 - Superscalar
 - Out-of-order execution
 - Speculative execution
- Thread level parallelism
 - Hyperthreading technology (aka SMT)
 - Multicore
- Process level parallelism
 - Multiprocessor system
 - Hyperthreading technology (aka SMT)
 - Multicore

That Was The Software

How about the Hardware?







Latency Vs Throughput

A Quick Glimpse on: Flynn Classification


- A taxonomy of computer architecture
- Proposed by Michael Flynn in 1966
- It is based on two things:
 - Instructions
 - Data

	Single instruction	Multiple instruction
Single data	SISD	MISD
Multiple data	SIMD	MIMD




PU = Processing Unit


More About MIMD




Shared Memory



Distributed Memory



Hybrid



Multicore and Manycore

We have arrived at many-core solutions not because of the success of our parallel software but because of our failure to keep increasing CPU frequency.*

Tim Mattson

Dilemma:

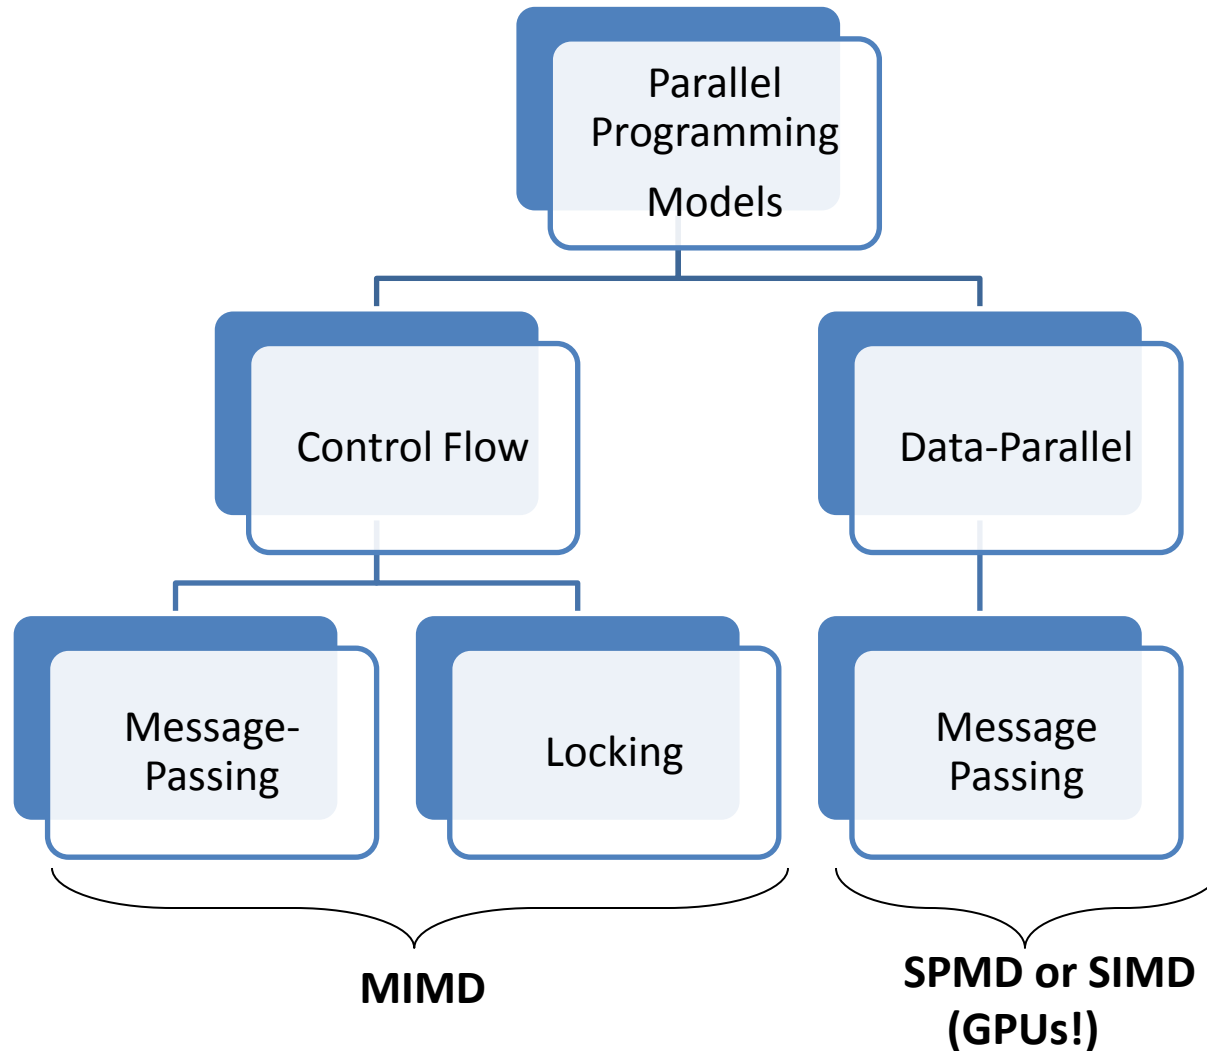
- Parallel hardware is ubiquitous
- Parallel software is not!

After more than 25 years of research, we are not closer to solving the parallel programming model!

The Mentality of Yet Another Programming Language ... Doesn't work!

ABCPL	CORRELATE	GLU	Mentat	Parafraze2	pC++
ACE	CPS	GUARD	Legion	Paralation	SCHEDULE
ACT++	CRL	HASL	Meta Chaos	Parallel-C++	SciTL
Active messages	CSP	Haskell	Midway	Parallaxis	SDDA
Adl	Cthreads	HPC++	Millipede	ParC	SHMEM
Adsmith	CUMULVS	JAVAR	CparPar	ParLib++	SIMPLE
ADDAP	DAGGER	HORUS	Mirage	ParLin	Sina
AFAPI	DAPPLE	HPC	MpC	Parmacs	SISAL
ALWAN	Data Parallel C	IMPACT	MOSIX	Parti	distributed smalltalk
AM	DC++	ISIS	Modula-P	pC	SMI
AMDC	DCE++	JAVAR	Modula-3*	PCN	SONiC
AppLeS	DDD	JADE	Multipol	PCP	Split-C
Amoeba	DICE	Java RMI	MPI	PH	SR
ARTS	DIPC	javaPG	MPC++	PEACE	Sdthreads
Athapascan-Ob	DOLIB	JavaSpace	Mumin	PCU	Strand
Aurora	DOME	JIDL	Nano-Threads	PET	SUIF
Automap	DOSMOS	Joyce	NESL	PENNY	Synergy
bb_threads	DRL	Khoros	NetClasses++	Phosphorus	Telegrphos
Blaze	DSM-Threads	Karma	Nexus	POET	SuperPascal
BSP	Ease	KOAN/Fortran-S	Nimrod	Polaris	TCGMSG
BlockComm	ECO	LAM	NOW	POOMA	Threads.h++
C*	Eiffel	Lilac	Objective Linda	POOL-T	TreadMarks
"C* in C	Eilem	Linda	Occam	PRESTO	TRAPPER
C++	Emerald	JADA	Omega	P-RIO	uC++
CarlOS	EPL	WWWinda	OpenMP	Prospero	UNITY
Cashmere	Excalibur	ISETL-Linda	Orca	Proteus	UC
C4	Express	ParLin	OOF90	QPC++	V
CC++	Falcon	Eilem	P++	PVM	ViC*
Chu	Filaments	P4-Linda	P3L	PSI	Visifold V-NUS
Charlotte	FM	POSYBL	Pablo	PSDM	VPE
Charm	FLASH	Objective-Linda	PADE	Quake	Win32 threads
Charm++	The FORCE	LiPS	PADRE	Quark	WinPar
Cid	Fork	Locust	Panda	Quick Threads	XENOOPS
Cilk	Fortran-M	Lparx	Papers	Sage++	XPC
CM-Fortran	FX	Lucid	AFAPI	SCANDAL	Zounds
Converse	GA	Maisie	Para++	SAM	ZPL
Code	GAMMA	Manifold	Paradigm		
COOL	Glenda				

Parallel Programming Models



Programming Model

- **Definition:** the languages and libraries that create an abstract view of the machine
- Control
 - How is parallelism created?
 - How are **dependencies** enforced?
- Data
 - Shared or private?
 - How is shared data accessed or private data communicated?
- Synchronization
 - What operations can be used to coordinate parallelism
 - What are the atomic (indivisible) operations?


It Is Important to Note

- You can run any paradigm on any hardware (e.g. an MPI on shared-memory)
- The same program can have different type of parallel paradigms
- The hardware itself can be heterogeneous

The whole challenge of parallel programming is to make the best use of the underlying hardware to exploit the different type of parallelisms

Example

We have a matrix A . We need to form another matrix $Asqr$ that contains the square of each element of A . Then we need to calculate S , which is the sum of the elements in $Asqr$.





- How can we parallelize this?
- How long will it take if we have unlimited number of processors?

Example

- First, decompose your problem into a set of tasks
 - There are many ways of doing it.
 - Tasks can be of the same, different, or undetermined sizes.
- Draw a task-dependency graph (do you remember the DAG we saw earlier?)
 - A directed graph with **Nodes** corresponding to tasks
 - **Edges** indicating dependencies, that the result of one task is required for processing the next.

Example



slide derived from Katherine Yelick

Does your knowledge of the
underlying hardware change
your task dependency graph?
If yes, how?

Conclusions

- Concurrency and parallelism are not exactly the same thing.
- There is parallelism at different granularities, with methods to exploit each parallelism granularity.
- You need to know the difference among: threads/processors/tasks.
- Knowing the hardware will help you generating a better task dependency graph.