



CSCI-GA.3033-012
**Multicore Processors:
Architecture & Programming**

**Lecture 7: OpenMP:
Control vs. Simplicity Tradeoff**

Mohamed Zahran (aka Z)

mzahran@cs.nyu.edu

<http://www.mzahran.com>



Small and Easy Motivation

```
int main() {  
  
    // Do this part in parallel  
  
    printf( "Hello, World!\n" );  
  
    return 0;  
}
```

Small and Easy Motivation

```
int main() {  
  
    omp_set_num_threads(16);  
  
    // Do this part in parallel  
    #pragma omp parallel  
    {  
        printf( "Hello, World!\n" );  
    }  
  
    return 0;  
}
```

Simple!

Serial Program:

```
void main()
{
    double Res[1000];

    for(int i=0;i<1000;i++) {
        do_huge_comp(Res[i]);
    }
}
```



Parallel Program:

```
void main()
{
    double Res[1000];
    #pragma omp parallel for
    for(int i=0;i<1000;i++) {
        do_huge_comp(Res[i]);
    }
}
```

OpenMP can parallelize many serial programs with relatively few annotations that specify parallelism and independence

OpenMP is a small API that hides cumbersome threading calls with simpler *directives*

Interesting Insights About OpenMP




These insights are coming from HPC folks though!

OpenMP In a Nutshell


- API
- Multithreaded programming
- Assumes shared memory multiprocessor
- Works with C/C++, and Fortran
 - example: `gcc -fopenmp`
- Consists of:
 - compiler directives
 - library routines
 - environment variables
- Strengths: portability, simplicity, and scalability
- Weakness: less control to the programmer

Goals of OpenMP

- Standardization among platforms/architectures
- Easy of use
- Portability



OpenMP uses the fork-join model of parallel execution.




All OpenMP programs begin with a single thread: **master thread** (ID = 0)

FORK: the master thread then creates a team of parallel *threads*.

JOIN: When the team threads complete the statements in the parallel region construct, they synchronize and terminate

Isn't Nested Parallelism Interesting?



Important!

- The following are implementation dependent:
 - Nested parallelism
 - Dynamically alter number of threads
- It is entirely up to the programmer to ensure that I/O is conducted correctly within the context of a multithreaded program.
- Threads can "cache" their data and are not required to maintain exact **consistency** with real memory all of the time. The programmer is responsible for insuring that the variable is FLUSHed by all threads as needed.

OpenMP

```
graph TD; OpenMP --> Directives; OpenMP --> RuntimeLibraries[Runtime Libraries]; OpenMP --> EnvirVariables[Envir. Variables]; Directives --> CodeExample["#pragma omp [clause, clause, ...]"];
```

Directives

Runtime Libraries

**Envir.
Variables**

`#pragma omp [clause, clause, ...]`

- Case sensitive
- Only one directive-name may be specified per directive
- Each directive applies to at most one succeeding statement, which must be a structured block.
- Long directive lines can be "continued" on succeeding lines by escaping the newline character with a backslash ("\") at the end of a directive line.

OpenMP

Directives

Runtime Libraries

**Envir.
Variables**

#pragma omp [clause, clause, ...]

Example:

```
#pragma omp parallel [clause ...] newline  
    if (scalar_expression)  
    private (list)  
    shared (list)  
    default (shared | none)  
    firstprivate (list)  
    reduction (operator: list)  
    copyin (list)  
    num_threads (integer-expression)
```

structured_block

```
#include <omp.h>
```

```
main () {
```

```
    int nthreads, tid;
```

```
    /* Fork a team of threads with each thread having a private tid variable */
```

```
    #pragma omp parallel private(tid) {
```

```
        /* Obtain and print thread id */
```

```
        tid = omp_get_thread_num();
```

```
        printf("Hello World from thread = %d\n", tid);
```

```
        /* Only master thread does this */
```

```
        if (tid == 0) {
```

```
            nthreads = omp_get_num_threads();
```

```
            printf("Number of threads = %d\n", nthreads);
```

```
        }
```

```
    }
```

```
    /* All threads join master thread and terminate */
```

```
}
```

```
#include <omp.h>
```

Runtime Libraries

```
main () {
```

```
    int nthreads, tid;
```

```
    /* Fork a team of threads with each thread having a private tid variable */
```

```
    #pragma omp parallel private(tid) {
```

```
        /* Obtain and print thread id */
```

```
        tid = omp_get_thread_num();
```

```
        printf("Hello World from thread = %d\n", tid);
```

```
        /* Only master thread does this */
```

```
        if (tid == 0) {
```

```
            nthreads = omp_get_num_threads();
```

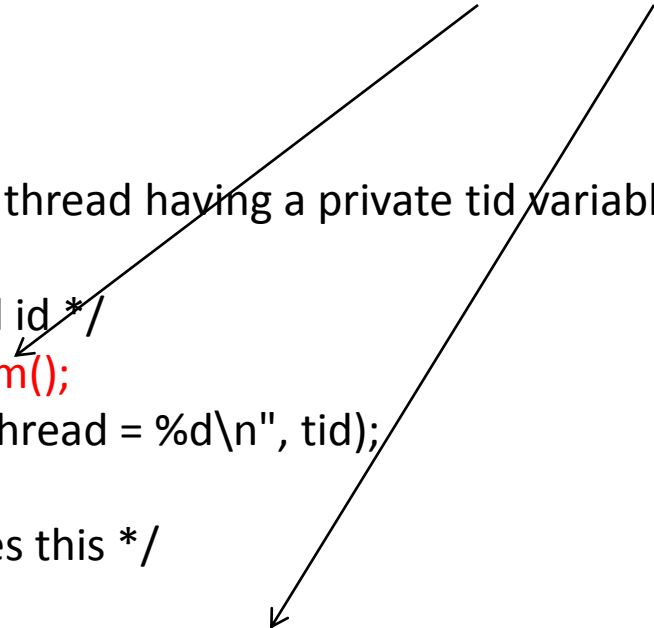
```
            printf("Number of threads = %d\n", nthreads);
```

```
        }
```

```
    }
```

```
    /* All threads join master thread and terminate */
```


```
}
```



How Many Threads?


- Setting of the **NUM_THREADS** clause
- Use of the **omp_set_num_threads()** library function
- Setting of the **OMP_NUM_THREADS** environment variable
- Implementation default - usually the number of cores.
- Threads are numbered from 0 (master thread) to N-1

Dividing Work Among Threads



Important: Work sharing directives do NOT launch new threads.

Dividing Work Among Threads



```
#pragma omp for [clause ...]  
                schedule (type [chunk])  
                ordered private (list)  
                firstprivate (list)  
                lastprivate (list)  
                shared (list)  
                reduction (operator: list)  
                collapse (n)  
                nowait
```

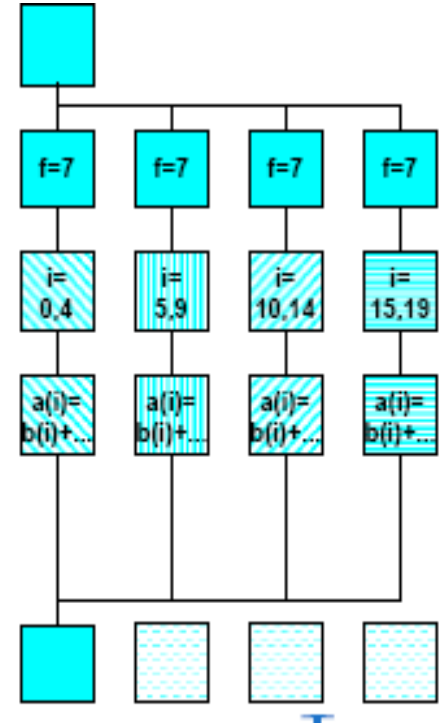
for_loop

- Number of iterations must be known in advance
- No gotos into or out of the loop.


```

#pragma omp parallel private(f)
{
    f=7;
    #pragma omp for
    for (i=0; i<20; i++)
        a[i] = b[i] + f * (i+1);
} /* omp end parallel */

```



Dividing Work Among Threads



```
#include <omp.h>
#define CHUNKSIZE 100
#define N 1000


main () {
    int i, chunk;
    float a[N], b[N], c[N];

    for (i=0; i < N; i++)
        a[i] = b[i] = i * 1.0;

    chunk = CHUNKSIZE;

    #pragma omp parallel shared(a,b,c,chunk) private(i)
    {
        #pragma omp for schedule(dynamic,chunk) nowait
        for (i=0; i < N; i++) c[i] = a[i] + b[i];
    } /* end of parallel section */
}
```

Dividing Work Among Threads



```
#pragma omp sections [clause ...]
```

```
{  
  #pragma omp section
```

```
    structured_block
```

```
  #pragma omp section
```

```
    structured_block
```


```
}
```

Implicit barrier here,
unless you use
nowait.

```


#pragma omp parallel
{
#pragma omp sections
  { a=...;
    b=...; }
#pragma omp section
  { c=...;
    d=...; }
#pragma omp section
  { e=...;
    f=...; }
#pragma omp section
  { g=...;
    h=...; }
} /*omp end sections*/
} /*omp end parallel*/

```



Dividing Work Among Threads

Specifies that the enclosed code is to be executed by **only one thread** in the team.



```
#pragma omp single [clause ...]
```

```
structured_block
```

Example About Consistency Model

Code:

Initially $A = \text{Flag} = 0$

P1

A = 23;
Flag = 1;

P2

while (Flag != 1) {;}
... = A;

Possible execution sequence on each processor:

P1

Write A 23
Write Flag 1

P2

Read Flag //get 0

.....

Read Flag //get 1
Read A //what do you get?

Do You see the problem?

Example About Consistency Model

Code:

Initially $A = \text{Flag} = 0$

P1

$A = 23;$

flush;

$\text{Flag} = 1;$

P2

$\text{while } (\text{Flag} \neq 1) \{ ; \}$

$\dots = A;$

Execution:

- P1 writes data into A
- **Flush waits till write to A is completed**
- P1 then writes data to Flag
- Therefore, if P2 sees $\text{Flag} = 1$, it is guaranteed that it will read the correct value of A even if memory operations in P1 before flush and memory operations after flush are reordered by the hardware or compiler.

What Does OpenMP Say Here?

`#pragma omp flush (list)`

- Thread-visible variables are written back to memory at this point.
- A bit complicated scenario arises if two threads execute the flush at the same time with common variables between them.
- If you do not specify a list, then all variables will be flushed

Synchronization: Critical Directive

Enclosed code

- executed by all threads, but
- **restricted to only one thread at a time**

- C/C++:

```
#pragma omp critical [ ( name ) ]
```

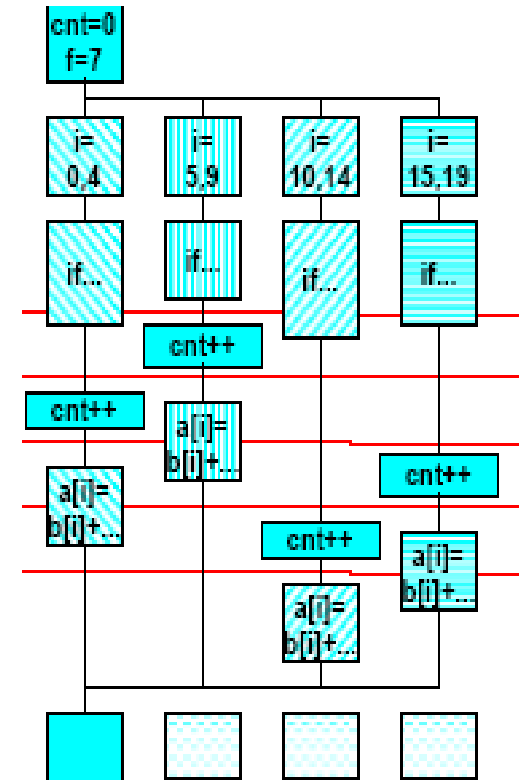
structured-block

- A thread waits at the beginning of a critical region until no other thread in the team is executing a critical region with the same name.

```

cnt = 0;
f=7;
#pragma omp parallel
{
  #pragma omp for
    for (i=0; i<20; i++) {
      if (b[i] == 0) {
        #pragma omp critical
          cnt ++;
      } /* endif */
      a[i] = b[i] + f * (i+1);
    } /* end for */
  } /*omp end parallel */

```



OpenMP



```
graph TD; OpenMP --> Directives; OpenMP --> RuntimeLibraries[Runtime Libraries]; OpenMP --> EnvirVariables[Envir. Variables];
```

Directives

Runtime Libraries

Envir. Variables

omp_get_thread_num()

- The number of threads remains unchanged for a parallel region
- The above function returns an integer
- Master thread has ID 0
- Different parallel regions may have different number of threads.

omp_get_num_threads()

- Returns, as integer, the number of threads in the current parallel region.

OpenMP



```
graph TD; OpenMP --> Directives; OpenMP --> RuntimeLibraries[Runtime Libraries]; OpenMP --> EnvirVariables[Envir. Variables];
```

Directives

Are we in a parallel region?

`OMP_IN_PARALLEL()`

How many processors in the system?

`OMP_GET_NUM_PROCS()`

Runtime Libraries

Envir. Variables

OpenMP



```
graph TD; OpenMP --> Directives; OpenMP --> RuntimeLibraries[Runtime Libraries]; OpenMP --> EnvirVariables[Envir. Variables];
```

Directives

Runtime Libraries

Envir. Variables

- Environment variables allow the end-user to control the parallel code.
- All environment variable names are uppercase.

Example:

OMP_NUM_THREADS

Sets the maximum number of threads to use during execution.

For example: **setenv OMP_NUM_THREADS 8**

What OpenMP Does NOT Do

- **Not Automatic parallelization**
 - User explicitly specifies parallel execution
 - Compiler does not ignore user directives even if wrong
- **Not meant for distributed memory parallel systems**
- **Not necessarily implemented identically by all vendors**
- **Not Guaranteed to make the most efficient use of shared memory**

Example:

```
#include <omp.h>

double MA[100][100], MB[100][100], MC[100][100];
int i, row, col, size = 100;

int main() {
    read_input(MA, MB);
    #pragma omp parallel shared(MA,MB,MC,size) private(row,col,i)
    {
        #pragma omp for schedule(static)
        for (row = 0; row < size; row++) {
            for (col = 0; col < size; col++)
                MC[row][col] = 0.0;
        }
        #pragma omp for schedule(static)
        for (row = 0; row < size; row++) {
            for (col = 0; col < size; col++)
                for (i = 0; i < size; i++)
                    MC[row][col] += MA[row][i] * MB[i][col];
        }
    }
    write_output(MC);
}
```


Example:

```
#include <omp.h>

double MA[100][100], MB[100][100], MC[100][100];
int i, row, col, size = 100;

int main() {
    read_input(MA, MB);
    #pragma omp parallel private(row,col,i)
    {
        #pragma omp for schedule(static)
        for (row = 0; row < size; row++) {
            #pragma omp parallel shared(MA, MB, MC, size)
            {
                #pragma omp for schedule(static)
                for (col = 0; col < size; col++) {
                    MC[row][col] = 0.0;
                    for (i = 0; i < size; i++)
                        MC[row][col] += MA[row][i] * MB[i][col];
                }
            }
        }
    }
    write_output(MC);
}
```

Questions: Pthreads Vs OpenMP

- When will you use Pthreads and when will you use OpenMP?
- Writing the same program in Pthreads and OpenMP, which one do you think will have better performance? Scalability?
- If you are using OpenMP, do you have any freedom to help the underlying hardware?

Conclusions

- OpenMP is an easier way than Pthreads for multithreaded programming
- OpenMP depends on compiler directives, runtime library, and environment variable.
- Many aspects of OpenMP are still implementation dependent, so you need to be careful!