



CSCI-GA.3033-012  
**Multicore Processors:**  
**Architecture & Programming**

**Lecture 12: Putting It All Together**

Mohamed Zahran (aka Z)

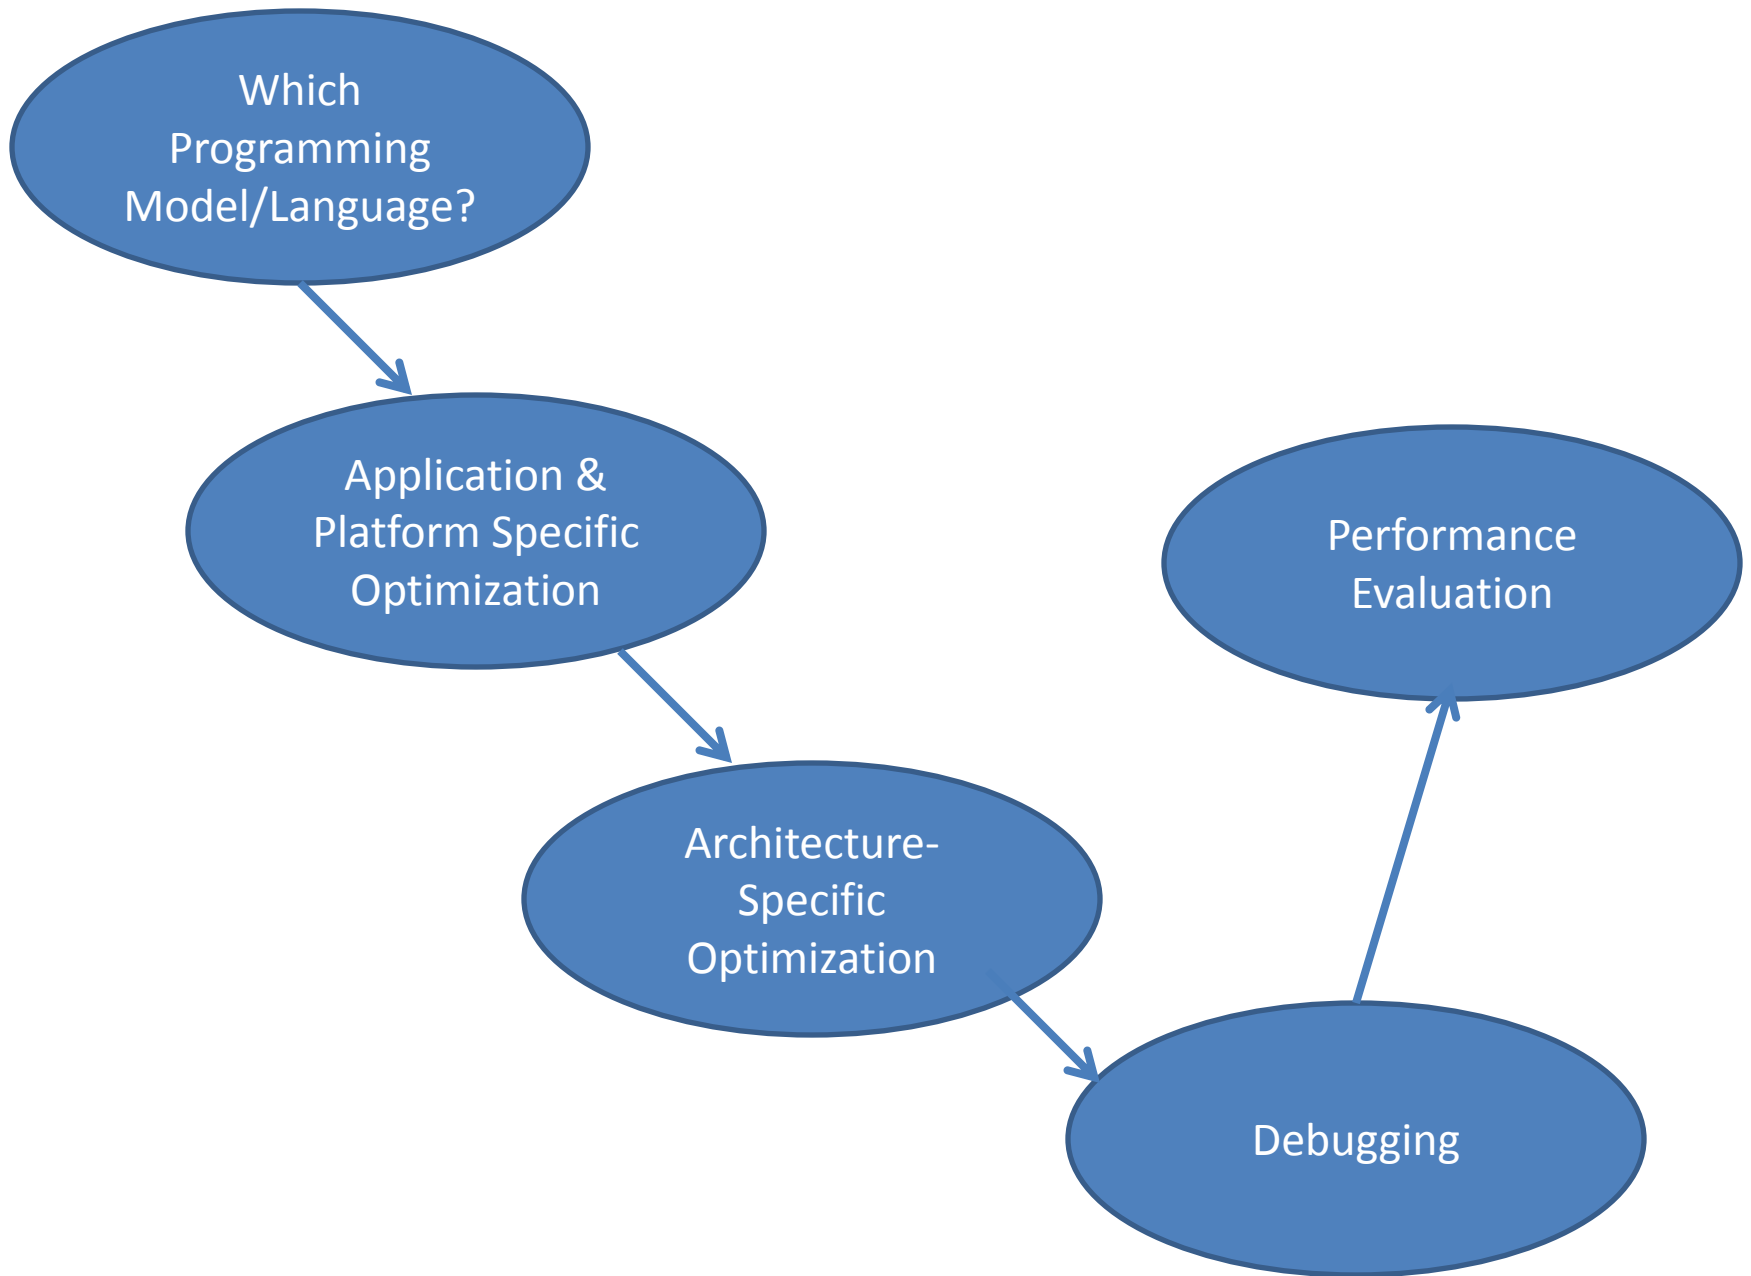
mzahran@cs.nyu.edu


<http://www.mzahran.com>



# Today's Lecture

- What to do with all what you heard about in this course?
- Challenges in hardware and how they affect the software
- Different type of parallelism
- Common problems in parallel programs
- A glimpse into the future
  - Software
  - Hardware





How to choose your programming language/model given an application?

# How to Choose?

- Your level of expertise
  - Less experience = higher abstraction
- The concurrency in the application
  - What type of parallelism we have?
- What can you do with the hardware?
- How does the platform interact with you?
- Are you starting from scratch? or from a sequential version of the program?

# Today's Lecture

- What to do with all what you heard about in this course?
- Challenges in hardware and how they affect the software
- Different type of parallelism
- Common problems in parallel programs
- A glimpse into the future
  - Software
  - Hardware


# Challenges

```
graph TD; A[Challenges] --> B[Power-Aware Software?]; A --> C[Reliability-Aware Software]
```

Power-Aware Software?

Reliability-Aware  
Software

# Power




\* "New Microarchitecture Challenges in the Coming Generations of CMOS Process Technologies" – Fred Pollack, Intel Corp. Micro32 conference key note - 1999.




# Cooking Aware Computing





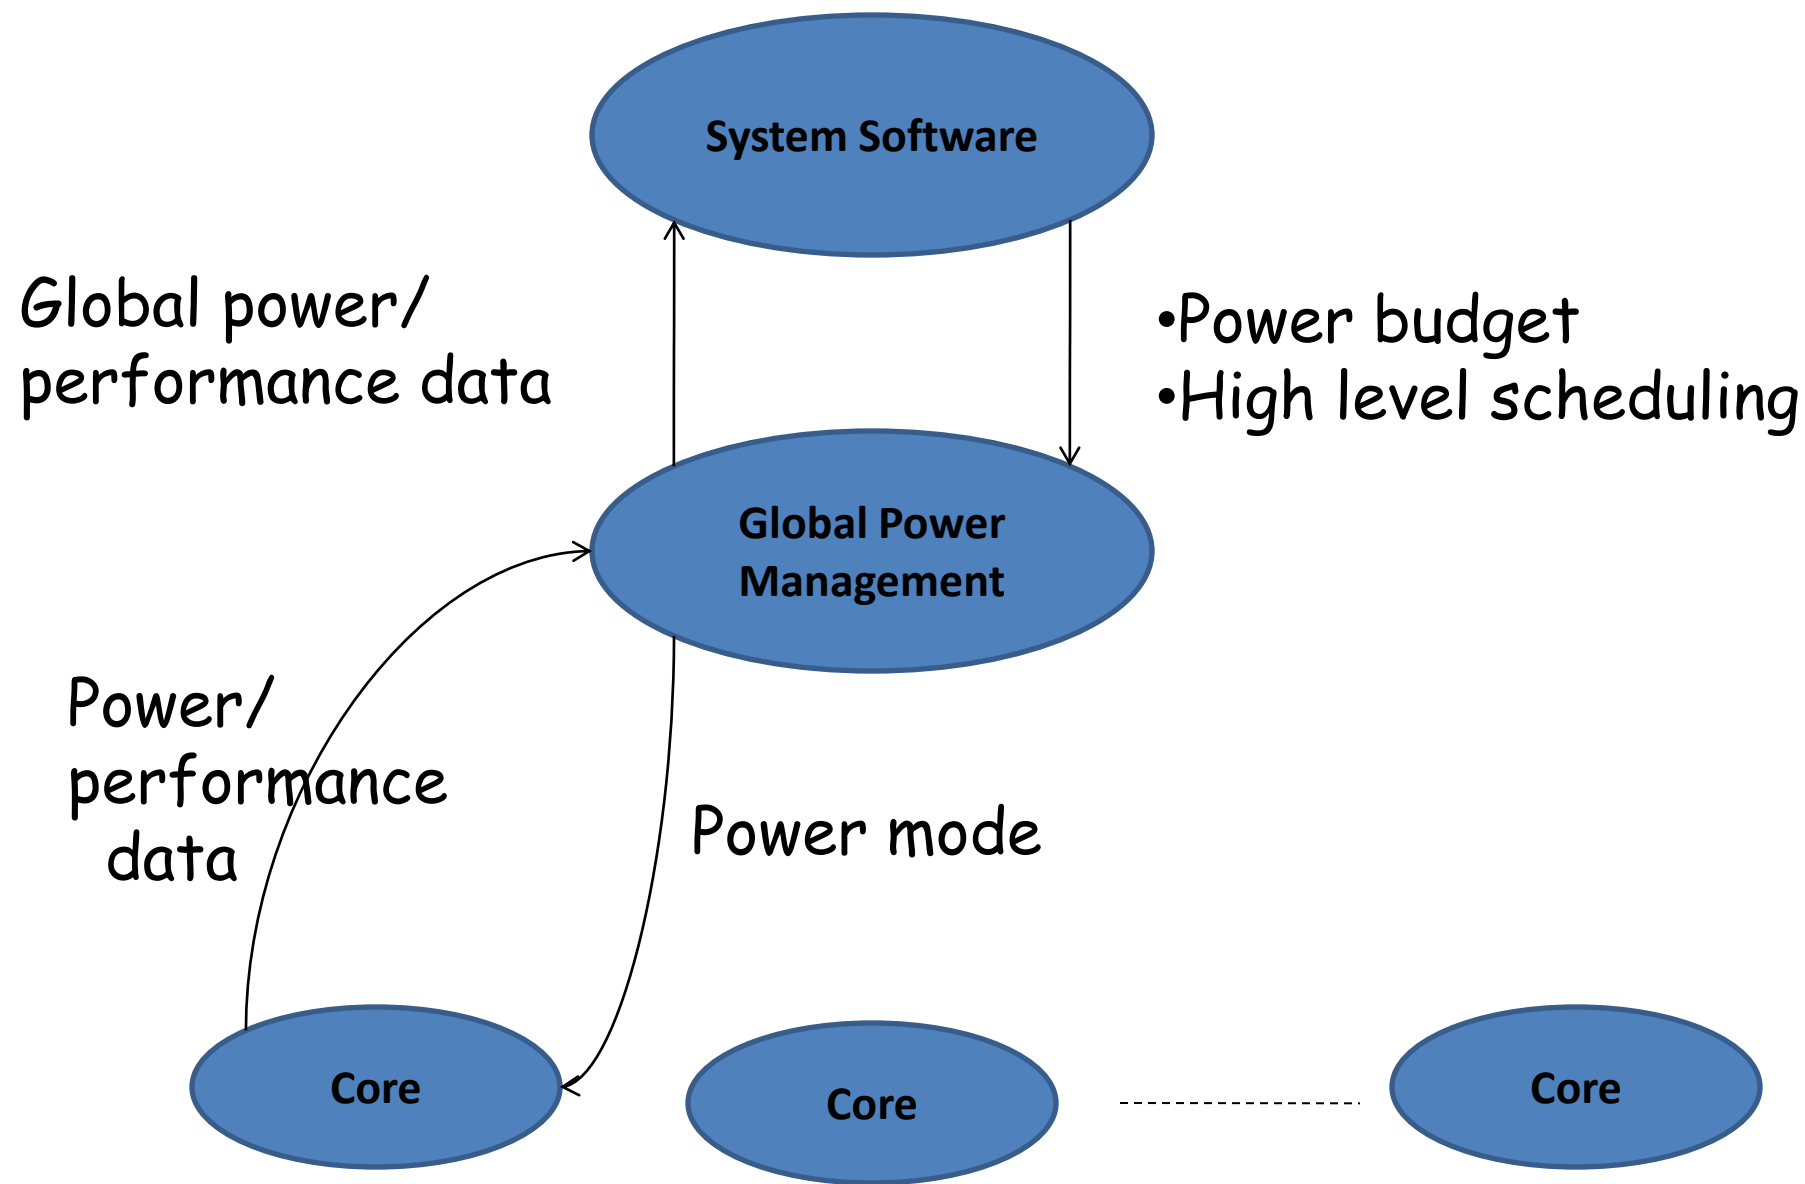
# Power-Aware Computing

- Definition: reducing power without losing performance
- Must deal with:
  - dynamic power
  - static power
  - temperature



# What To Do About Dynamic Power

- Stop and go
- DFVS (Dynamic Frequency and Voltage Scaling)
  - At OS level
    - idle time represents energy waste
    - deadlines for interactive programs
  - Offline compiler analysis
    - insert mode-set instructions
    - depends on program phases
    - lowers the voltage for memory-bound sections
  - Online dynamic compiler analysis
    - phase detection
    - binary instrumentation
- Reducing switching activity

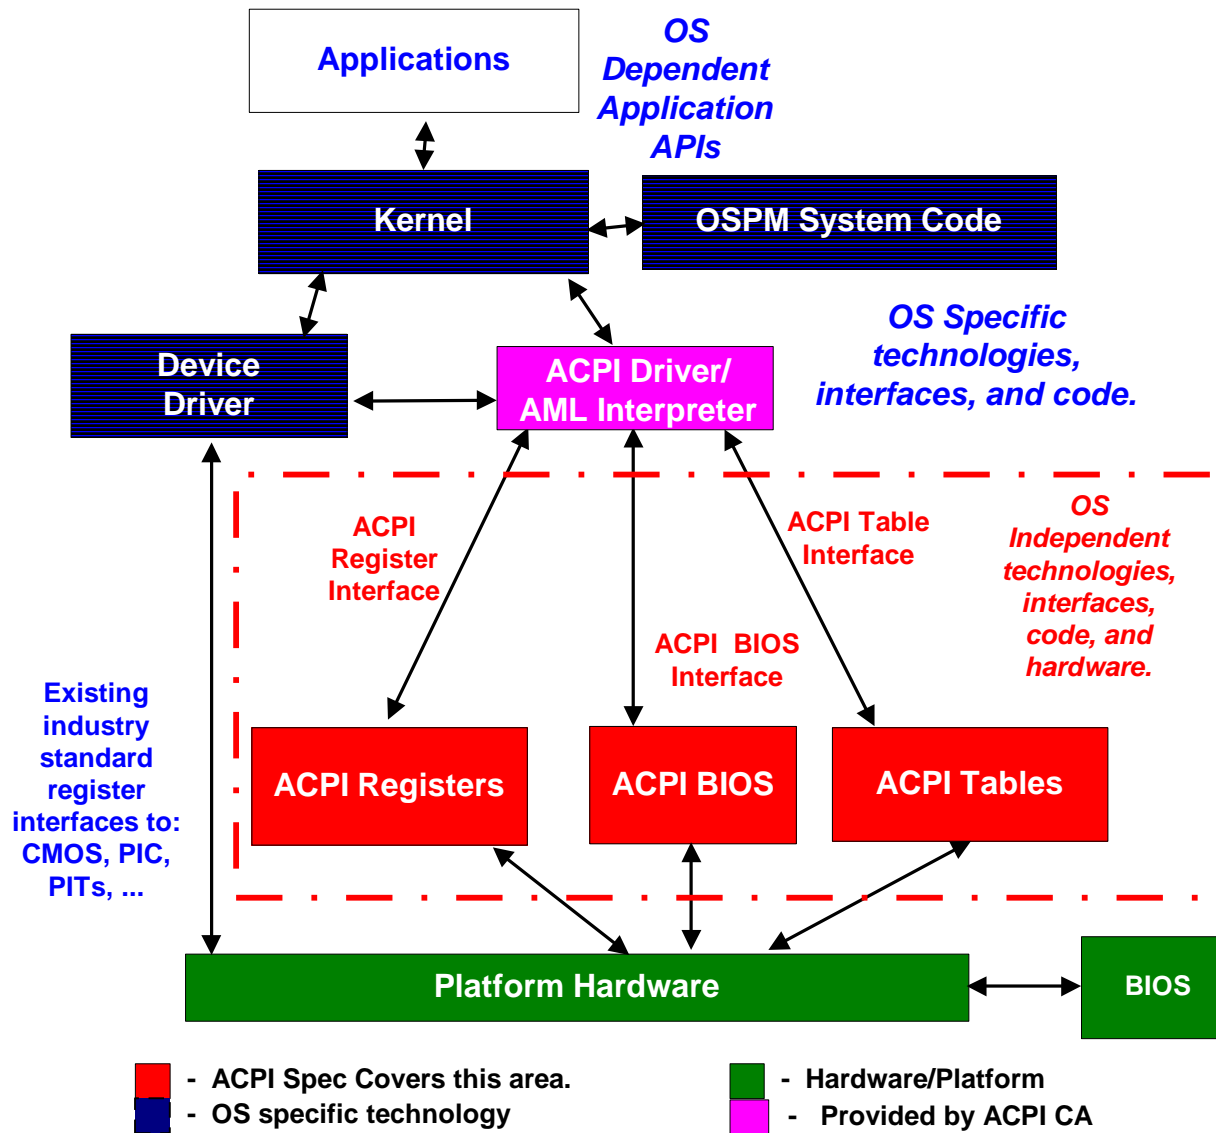


# Other Techniques for the Multicore

- Migration
  - Moving threads among cores
  - Timescale on order of a millisecond, much slower than DVFS
  - Migration can be used with DVFS
- Migrate critical thread
  - Measure criticality with heat sensor
  - Or with cache misses as a proxy

# As A Programmer

- Try to loose less-expensive operations (i.e. help the compiler)
- Locality to help caches
- Control DVFS
  - Advanced Configuration and Power Interface (ACPI)







# Reliability



- Transistors are becoming unreliable
- What will your application do if a core becomes unavailable?
- Can you duplicate some computations for very critical operations?

# Today's Lecture

- What to do with all what you heard about in this course?
- Challenges in hardware and how they affect the software
- Different type of parallelism
- Common problems in parallel programs
- A glimpse into the future
  - Software
  - Hardware

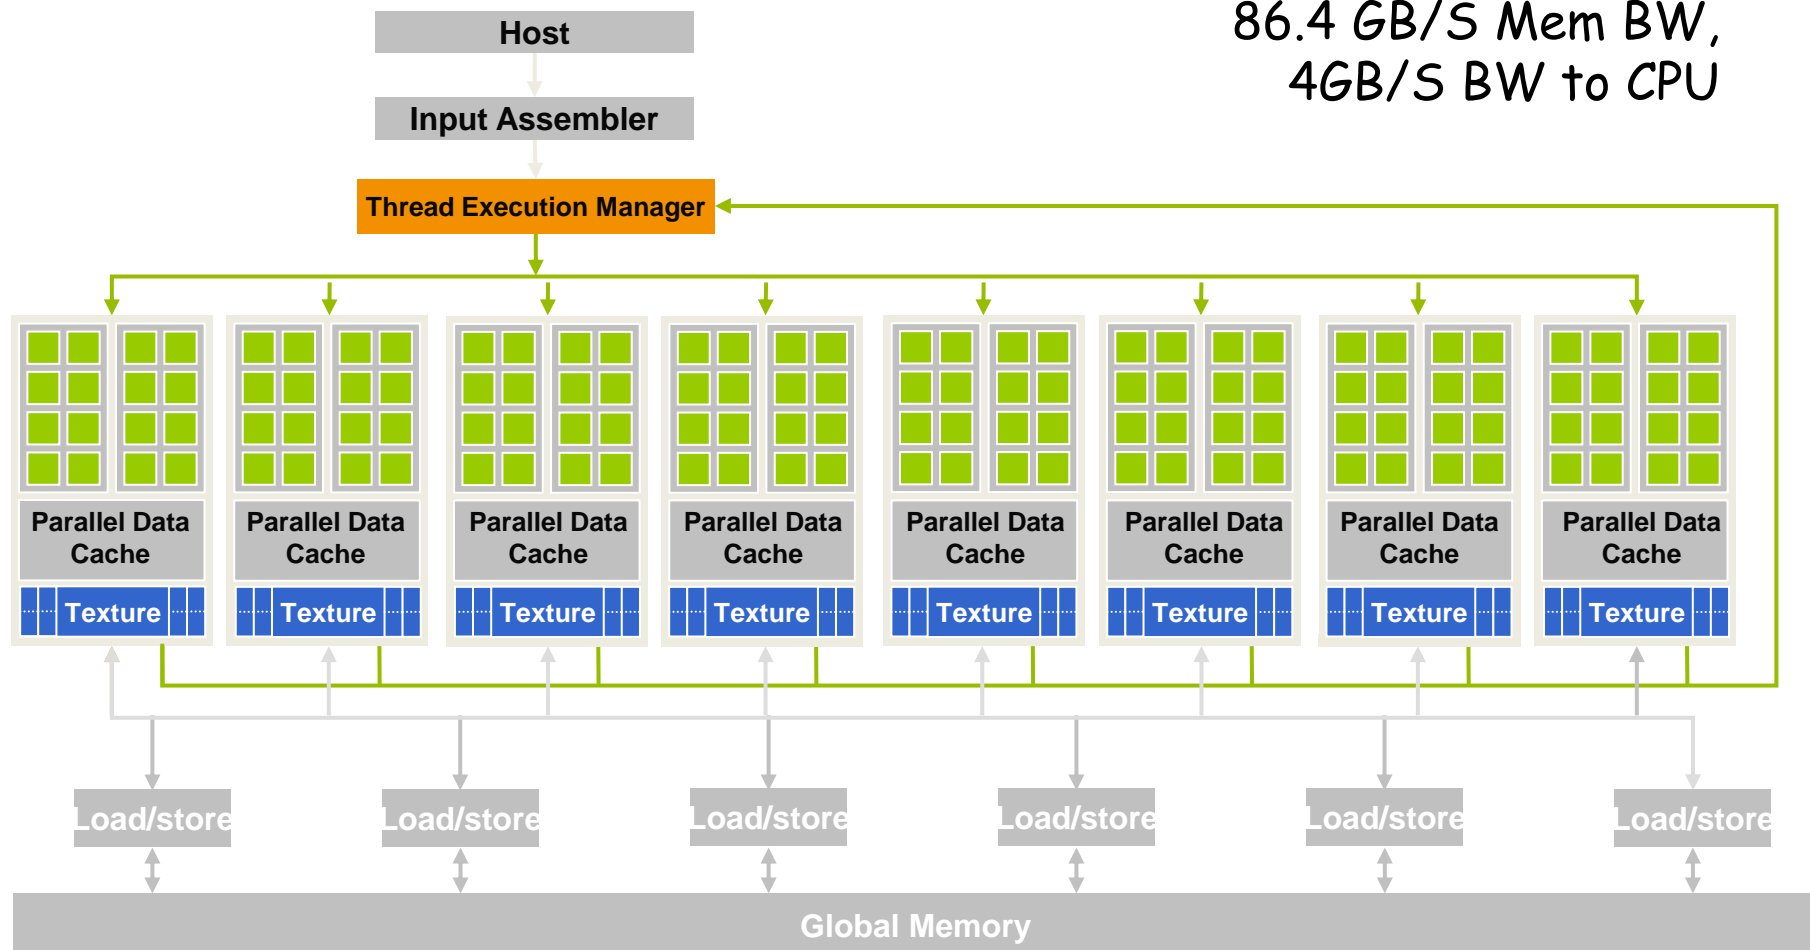
# A Glimpse at Another Type of Parallelism: SIMD/SPMD/STMD






# A Glimpse at A Typical GPU: GeForce 8800 (2007)

16 highly threaded SM's, >128 FPU's,  
367 GFLOPS, 768 MB DRAM,  
86.4 GB/S Mem BW,  
4GB/S BW to CPU



# A Glimpse at A Typical GPU

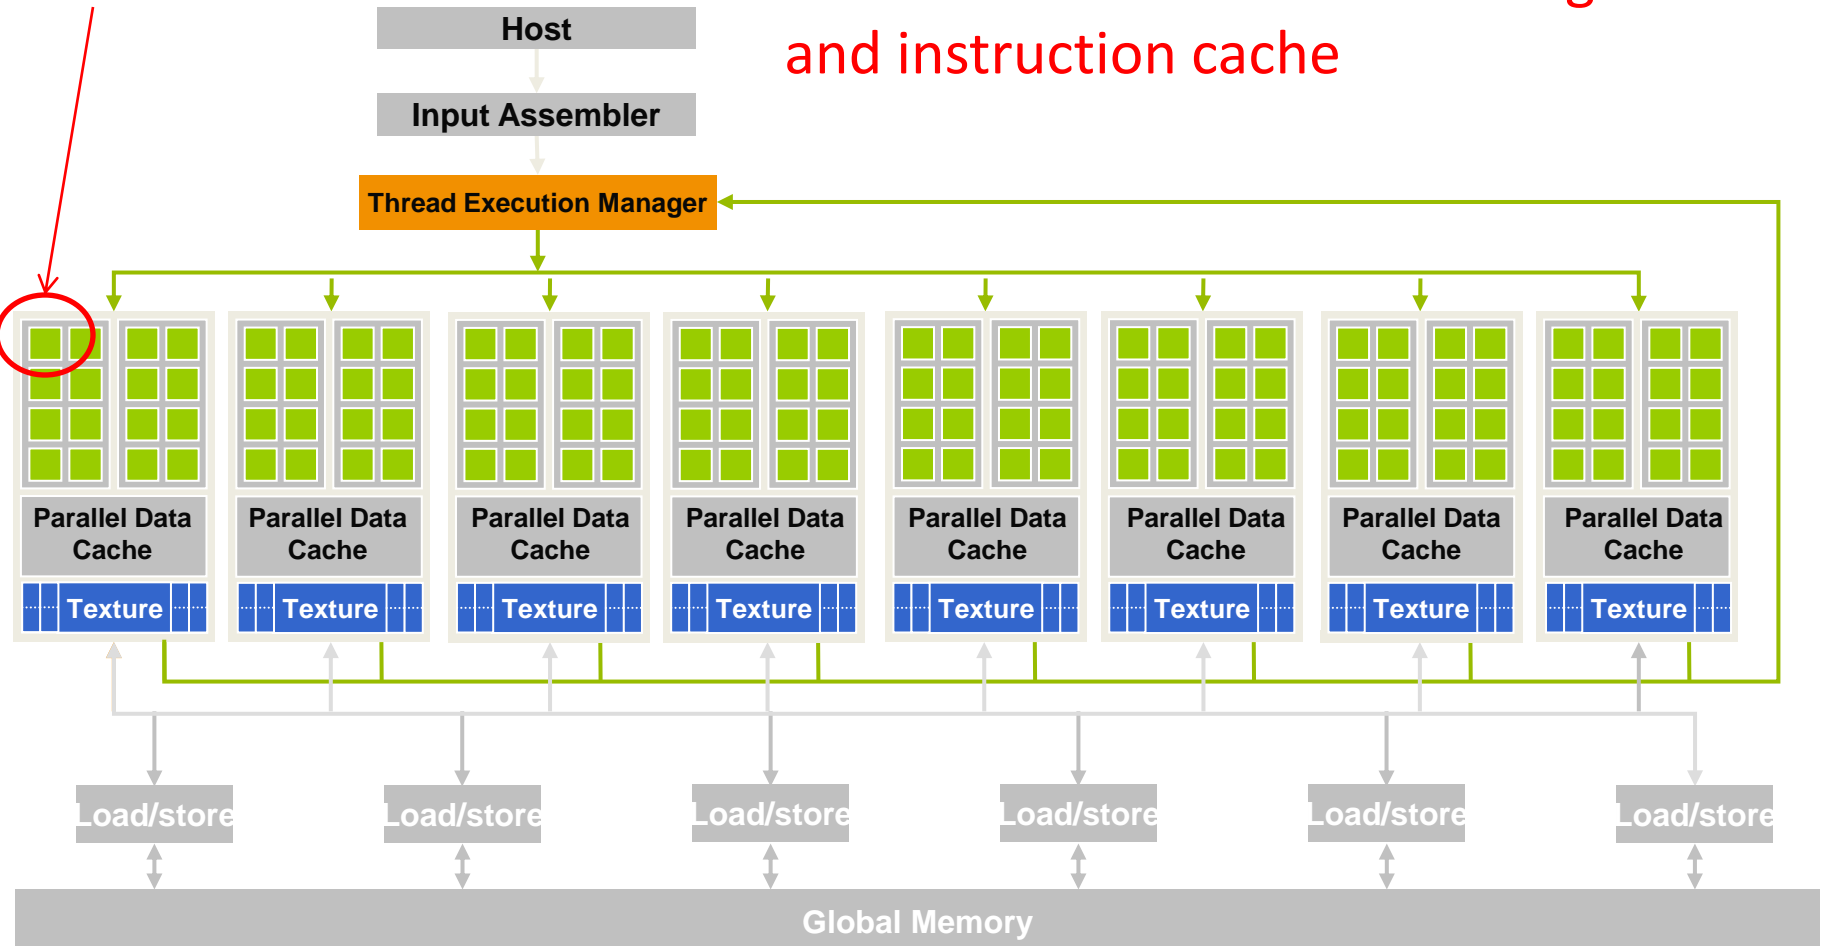
Streaming Multiprocessor (SM)



# A Glimpse at A Typical GPU

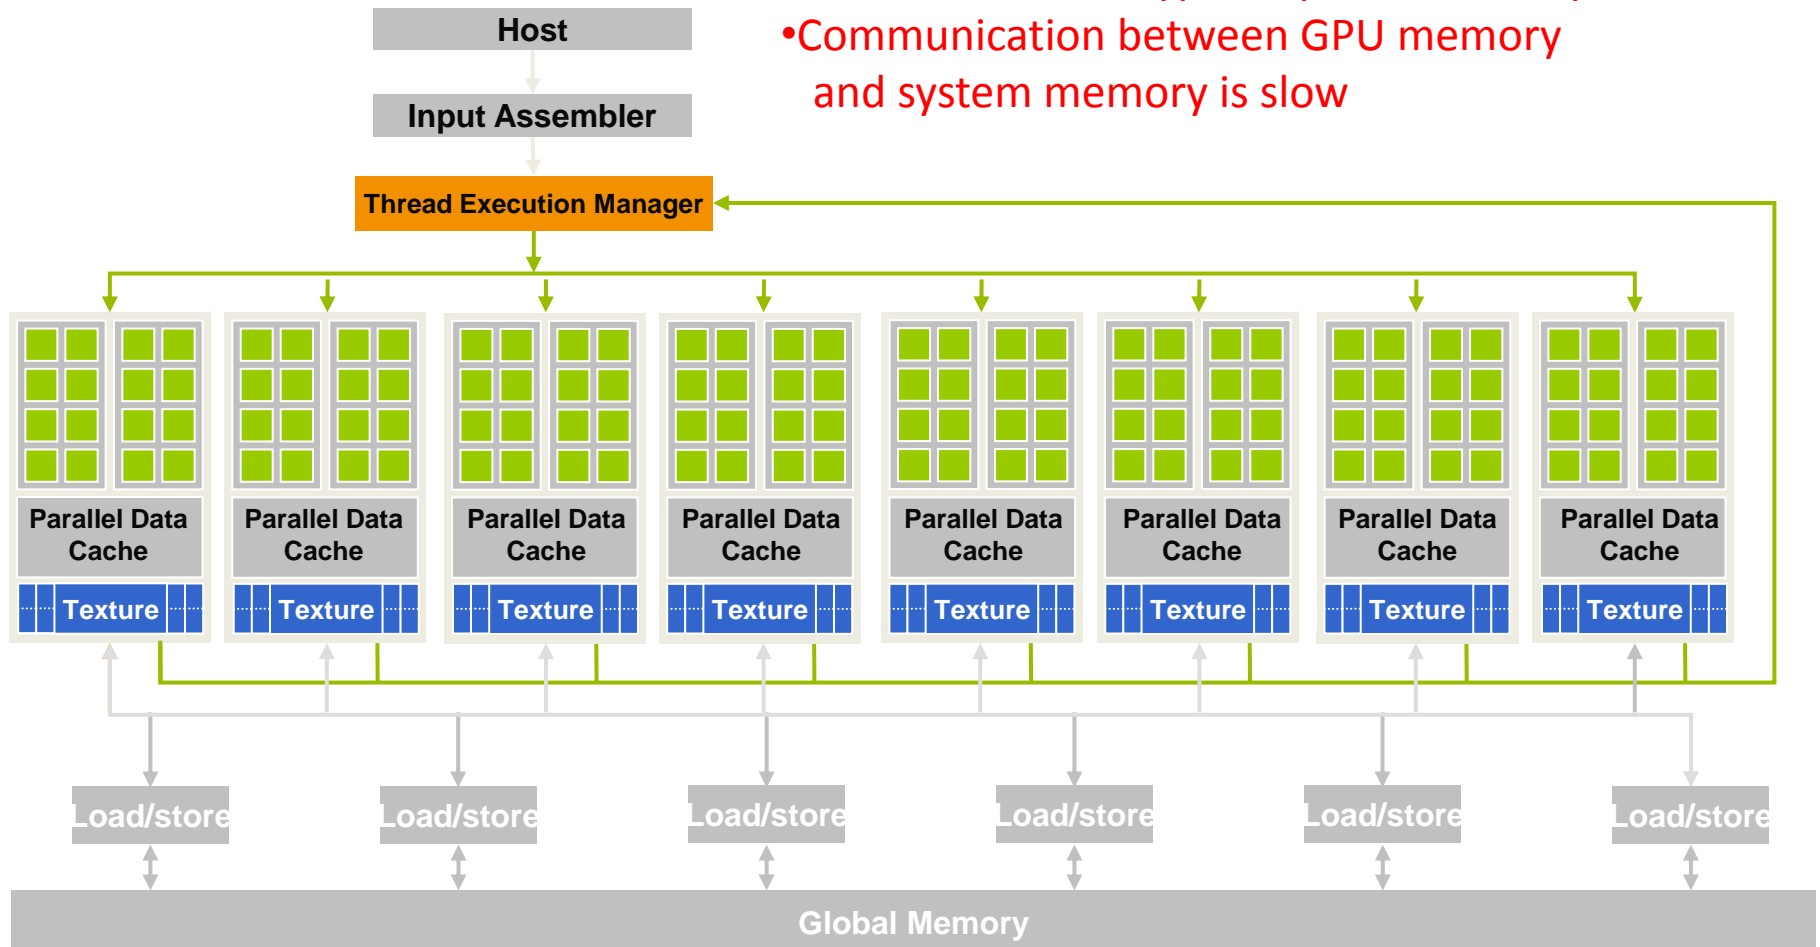
Streaming  
Processor (SP)

SPs within SM share control logic  
and instruction cache



# A Glimpse at A Typical GPU

- Much higher bandwidth than typical system memory
- A bit slower than typical system memory
- Communication between GPU memory and system memory is slow





# BUT ...

- GPU is not standalone, it needs CPU
- How to divide your program among multicore and GPU?
- Can you put part of your program into STMD/SPMD/SIMD?
- GPU → OpenCL and CUDA

# Today's Lecture

- What to do with all what you heard about in this course?
- Challenges in hardware and how they affect the software
- Different type of parallelism
- Common problems in parallel programs
- A glimpse into the future
  - Software
  - Hardware

# Common problems in parallel programs

- Parallel programs are subject to the usual bugs
- Plus: new timing and synchronization bugs (race condition, deadlocks, livelocks, ... )
- parallel bugs often disappear when you add code to try to identify the bug ☹️

# Common problems in parallel programs: Too Many Threads

- (Fixed amount of work)/(large number of threads) → each thread will do too little + the overhead of threading
- overhead of threading:
  - starting ending threads
  - contention on shared resources → especially when software threads are more than hardware threads

# Common problems in parallel programs: Too Many Threads

- The best strategy: limit the number of software threads to:
  - Number of hardware threads
  - Number of outer-level caches
- Also, separate your threads into I/O threads and compute threads
  - Blocked threads are not fighting for time-slice by the OS
- OpenMP takes this burden from the programmer
- Thread pool is another option (no standard one for POSIX threads though)

# Common problems in parallel programs: Data Race, Deadlocks, and Live Locks

- Sometimes races are hidden by the language syntax (What may seem like a single instruction may actually be several ones.)
- Several ways to deal with that:
  - Use tools like Intel Thread checker
  - Locks
  - Transactions (which may go down to locks in some implementations!).

# Common problems in parallel programs: Data Race, Deadlocks, and Live Locks

- Locks can lead to deadlocks
- Deadlocks occur when the following 4 conditions exist:
  - Access to each resource is exclusive
  - A thread is allowed to hold one resource while requesting another
  - No thread is willing to relinquish a resource it has acquired
  - There is a cycle of threads trying to acquire resources
- Deadlocks can be avoided by breaking anyone of the above four conditions, for example:
  - Replicate a resource if possible
  - Always let threads acquire locks (resources) in the order

# Common problems in parallel programs: Heavily Contended Locks

- Locks can become contended → performance degradation
- What to do?
  - Replicating the resource (hence spreading the contention) can help.
  - Consider partitioning the resource and use locks for each part



# Today's Lecture

- What to do with all what you heard about in this course?
- Challenges in hardware and how they affect the software
- Different type of parallelism
- Common problems in parallel programs
- A glimpse into the future
  - Software
  - Hardware



# How Will the Future Look Like?

- "I think there is a world market for maybe five computers."
  - Thomas Watson, chairman of IBM, 1949
- "There is no reason in the world anyone would want a computer in their home. No reason."
  - Ken Olsen, Chairman, DEC, 1977
- "640K of RAM ought to be enough for anybody."
  - Bill Gates, 1981

Predicting the Future is not easy!!

# How Will the Future Look Like?

- **Evolution:** just interpolating the current trend
- **Revolution:** a new technology, a paradigm shift → very hard to predict!

# The Future of Software: Evolution

- Application types
- Languages
- MPI for multicore?
- Auto-parallelization
- Functional-programming
- Transactional memories

# What Kind of Apps Need 100s Cores?

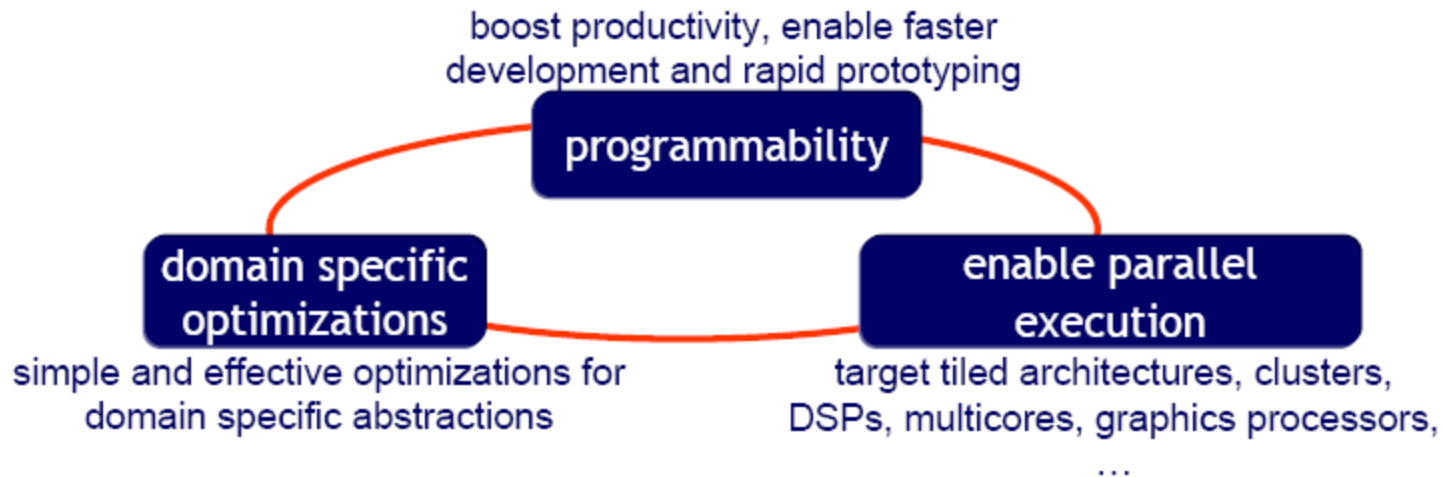
- “Who needs 100 cores to run M/S Word?”
  - Need compelling apps that use 100s of cores
- Compelling in terms of likely market or social impact, with longer term potential

# Example of Evolving Applications

- Content-based image retrieval
- Health-record management
- NLP
- Teleconferences
- Heavy multimedia contents
- More realistic graphics and user interface
- Event-driven (real-time)
- ...

# Evolving Languages

- Scripting
- Domain specific languages (e.g. StreamIt, ...)



# Message Passing Interface (MPI)

- Message Passing Model
- Allows communication between processes (threads) using specific message-passing system calls.
- Shared data is communicated through messages
- Does not assume shared memory



# Message Passing Interface (MPI)

- Allows for asynchronous events
- Does not require programmer to write in terms of loop-level parallelism
- Operates on multicore AND is scalable to very large systems
- Extremely flexible

**Why then it was not adopted for multicore as the de facto?**


# Message Passing Interface (MPI)

- Considered extremely difficult to write
  - Shared-memory models seem more intuitive
- Difficult to incrementally increase parallelism

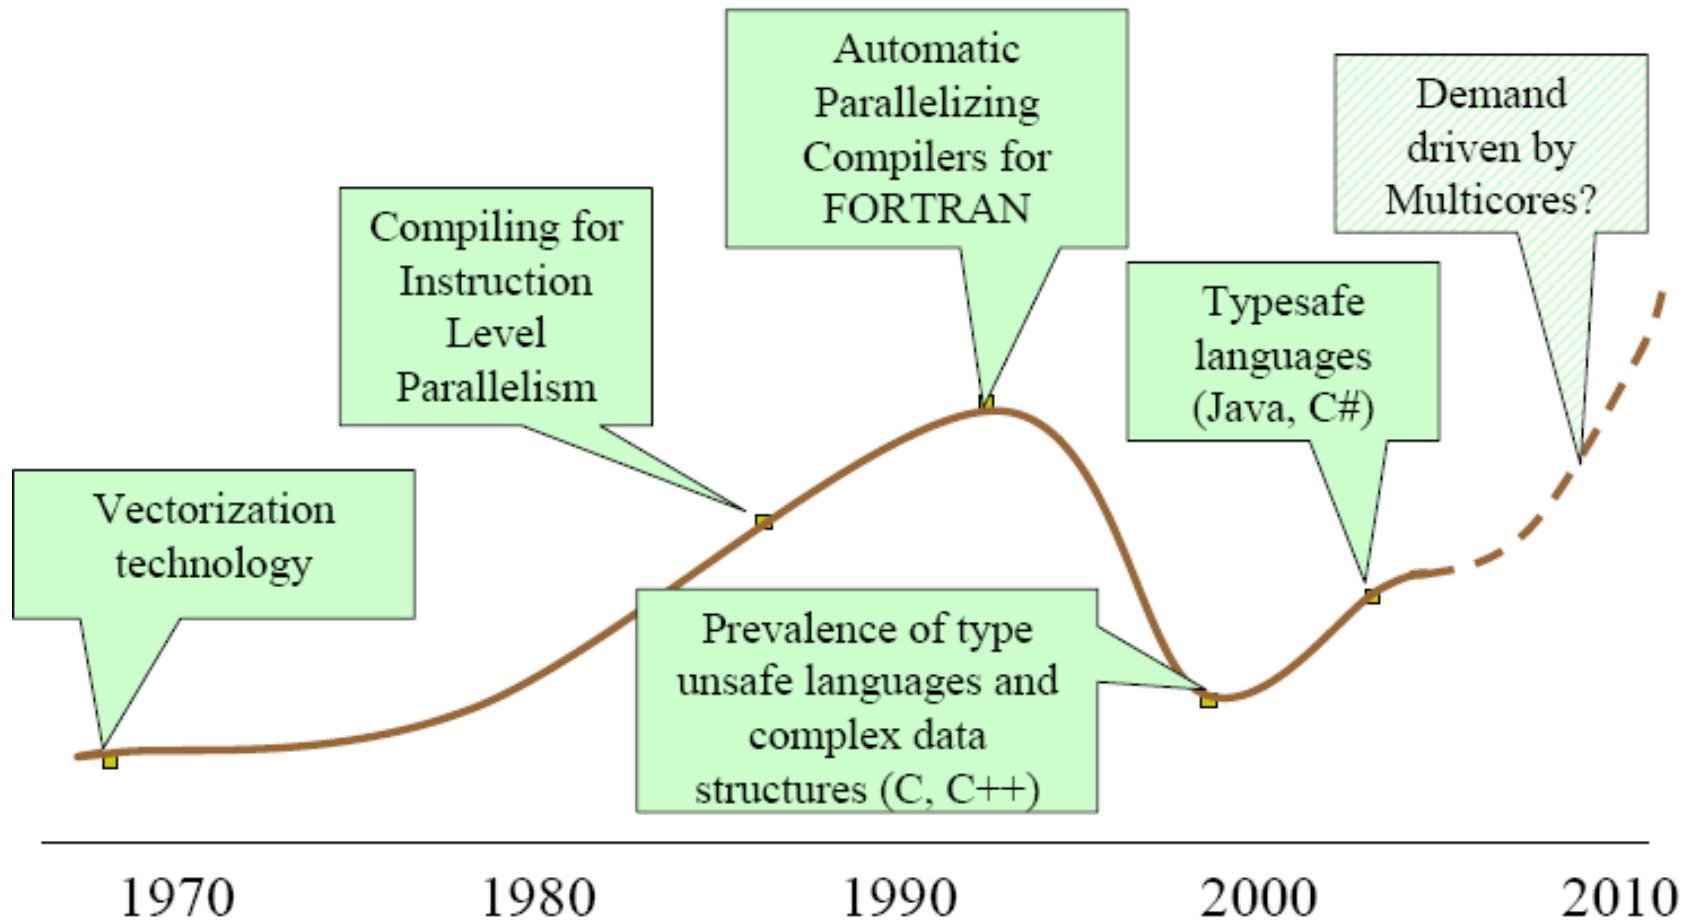
# MPI on Multicore

- One MPI process per core
  - Each MPI process is a single thread
- One MPI process per node
  - MPI processes are multithreaded
  - One thread per core
  - aka Hybrid model

# MPICH2 Features to Support Multicore Architecture



# Improvement in Automatic Parallelization



**Source:** Saman Amarasinghe, MIT

# What Do We Need From the Compiler?

- Compilers are critical in reducing the burden on programmers
  - Identification of data parallel loops can be easily automated, but many current systems require the programmer to do it.
- Reviving the push for automatic parallelization
  - Best case: totally automated parallelization hidden from the user
  - Worst case: simplify the task of the programmer

# The Future of Software: Revolution

- Program design methodology
  - Sketching?
- New programming paradigm

# The Future of Hardware: Evolution

- More cores on-chip
- Constraints increase
  - Dark-silicon becomes more substantial
  - Wire delay
  - Power
  - Reliability
- Clustered architecture on-chip
- More widespread heterogeneous multicore



# The Future of Hardware: Revolution

- Brain-inspired machines
- Reconfigurable processors
- DNA computing
- Quantum computing
- Non Von-Neumann model
- ...

# Conclusions

- Multicore and manycore processors is a work in progress → new techniques/developments almost daily
- No magical recipe → must keep track of the software and hardware

Thank You!