


# Pthreads: Busy Wait, Mutexes, Semaphores, Conditions, and Barriers

Bryan Mills

Chapter 4.2-4.8

# A Shared Memory System



# Pthreads

- Pthreads is a POSIX standard for describing a thread model, it specifies the API and the semantics of the calls.
- Model popular –practically all major thread libraries on Unix systems are Pthreads-compatible
- The Pthreads API is only available on POSIXR systems — Linux, MacOS X, Solaris, HPUX, ...

# Preliminaries


- Include `pthread.h` in the main file
- Compile program with `-lpthread`
  - `gcc -o test test.c -lpthread`
  - may not report compilation errors otherwise but calls will fail
- Check return values on common functions

# Include

```
#include <stdio.h>
```

```
#include <pthread.h>
```

This includes the  
pthreads library.



```
#define NUM_THREADS 10
```

```
void *hello (void *rank) {  
    printf("Hello Thread");  
}
```

```
int main() {  
    pthread_t ids[NUM_THREADS];  
    for (int i=0; i < NUM_THREADS; i++) {  
        pthread_create(&ids[i], NULL, hello, &i);  
    }  
    for (int i=0; i < NUM_THREADS; i++) {  
        pthread_join(ids[i], NULL);  
    }  
}
```

# C Libraries

- Libraries are created from many source files, and are either built as:
  - Archive files (libmine.a) that are statically linked.
  - Shared object files (libmine.so) that are dynamically linked
- Link libraries using the gcc command line options -L for the path to the library files and -l to link in a library (a .so or a .a):

```
gcc -o myprog myprog.c -L/home/bnm/lib -lm
```

- You may also need to specify the library path:  
-I /home/bnm/include

# Runtime Linker

- Libraries that are dynamically linked need to be able to find the shared object (.so)
- `LD_LIBRARY_PATH` environment variable is used to tell the system where to find these files.

```
export LD_LIBRARY_PATH=/home/bnm/lib:$LD_LIBRARY_PATH
```

# Header Files

- Header files (.h) contain function and macro definitions to be shared with other programs.
- To be used to specify the functions provided by the actual code files (.c, .cc)



# Thread creation

- Types: `pthread_t` – type of a thread

- Functions:

```
int pthread_create(pthread_t *thread,  
                  const pthread_attr_t *attr,  
                  void * (*start_routine)(void *),  
                  void *arg);  
  
int pthread_join(pthread_t thread, void **status);  
int pthread_detach();  
void pthread_exit();
```

- No explicit parent/child model, except main thread holds process info
- Call `pthread_exit` in main, don't just fall through;
- Most likely you wouldn't need `pthread_join`
  - `status` = exit value returned by joinable thread
- Detached threads are those which cannot be joined (can also set this at creation)

# pthread\_t objects

- Opaque
- Actual data that they store is system-specific.
- Data members aren't directly accessible to user code.
- However, the Pthreads standard guarantees that a pthread\_t object does store enough information to uniquely identify the thread with which it's associated.

# pthread\_t object

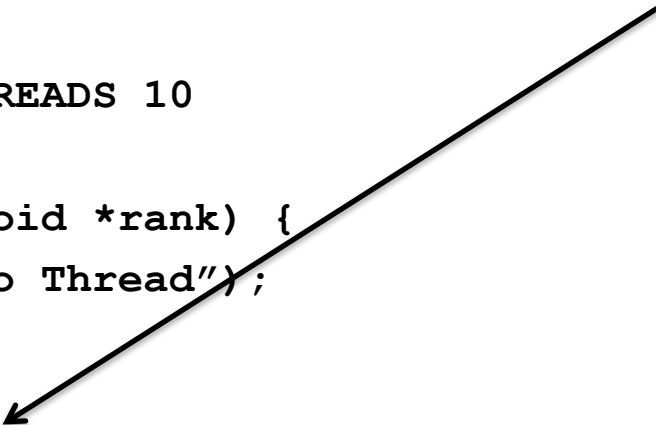
```
#include <stdio.h>
#include <pthread.h>

#define NUM_THREADS 10

void *hello (void *rank) {
    printf("Hello Thread");
}

int main() {
    pthread_t ids[NUM_THREADS];
    for (int i=0; i < NUM_THREADS; i++) {
        pthread_create(&ids[i], NULL, hello, &i);
    }
    for (int i=0; i < NUM_THREADS; i++) {
        pthread_join(ids[i], NULL);
    }
}
```

Creating an array of pthread\_t objects to hold all the threads we create.



# Create Thread

```
int pthread_create (
    pthread_t*  thread_p          /* out */ ,
    const pthread_attr_t* attr_p  /* in */ ,
    void*  (*start_func ) ( void ) /* in */ ,
    void*  arg_p                  /* in */ );
```

- thread\_p – Will set/return an allocated pthread\_t object with information about the created thread.
- attr\_p – Used to specify thread properties, NULL = Defaults.
- start\_func – Pointer to the function to execute.
- arg\_p – Argument to the start\_func.
- Returns 0 on success, non-zero on failure.

# Attributes

- **Type:** `pthread_attr_t` (see `pthread_create`)
- Attributes define the state of the new thread
- Attributes: system scope, joinable, stack size, inheritane.
- Use default behaviors with `NULL`.

```
int pthread_attr_init(pthread_attr_t *attr);  
int pthread_attr_destroy(pthread_attr_t *attr);  
pthread_attr_{set/get}{attribute}
```

- **Example:**

```
pthread_attr_t attr;  
pthread_attr_init(&attr); // Needed!!!  
pthread_setdetachstate(&attr, PTHREAD_CREATE_DETACHED);  
pthread_attr_setscope(&attr, PTHREAD_SCOPE_SYSTEM);  
pthread_create(NULL, &attr, foo, NULL);
```

# Function started by pthread\_create

- Prototype:  
`void* thread_function ( void* args_p ) ;`
- Void\* can be cast to any pointer type in C.
- So args\_p can point to a list containing one or more values needed by thread\_function.
- Similarly, the return value of thread\_function can point to a list of one or more values.

# pthread\_create

```
#include <stdio.h>
#include <pthread.h>
```

```
#define NUM_THREADS 10
```

```
void *hello (void *rank) {
    printf("Hello Thread");
}
```

```
int main() {
    pthread_t ids[NUM_THREADS];
    for (int i=0; i < NUM_THREADS; i++) {
        pthread_create(&ids[i], NULL, hello, &i);
    }
    for (int i=0; i < NUM_THREADS; i++) {
        pthread_join(ids[i], NULL);
    }
}
```

Creating new thread(s).

Keeping track of threads.

Default attributes.

Function to call.

Argument to pass to  
function to call.

# Stopping the Threads

- We call the function `pthread_join` once for each thread.
- A single call to `pthread_join` will wait for the thread associated with the `pthread_t` object to complete.



# pthread\_join

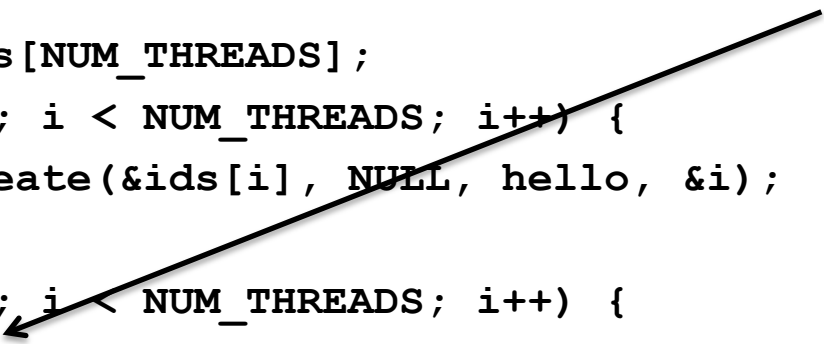
```
#include <stdio.h>
#include <pthread.h>
```

```
#define NUM_THREADS 10
```


```
void *hello (void *rank) {
    printf("Hello Thread");
}
```

```
int main() {
    pthread_t ids[NUM_THREADS];
    for (int i=0; i < NUM_THREADS; i++) {
        pthread_create(&ids[i], NULL, hello, &i);
    }
    for (int i=0; i < NUM_THREADS; i++) {
        pthread_join(ids[i], NULL);
    }
}
```

Called once for each thread. Will wait on thread to finish.



# Running the Threads




- Main thread forks and joins two threads.
  - Not really a parent->child relationship
  - Only distinction is that main thread holds system process information.

# Contention Scope

- *Contention scope* is the POSIX term for describing bound and unbound threads
- A bound thread is said to have *system contention scope*
  - i.e., it contends with all threads in the system
- An unbound thread has *process contention scope*
  - i.e., it contends with threads in the same process
- For most practical reasons use bound threads (system scope)
  - Solaris LWP switching cheap, Linux is one-to-one anyways...

# What's Wrong?

```
void *hello (void *rank) {  
    int *rank_int_ptr = (int*) rank;  
    printf("Hello Thread %d\n", *rank_int_ptr);  
}  
  
int main() {  
    pthread_t ids[NUM_THREADS];  
    for (int i=0; i < NUM_THREADS; i++) {  
        pthread_create(&ids[i], NULL, hello, &i);  
    }  
}
```



What `i` will be passed in?

# Global variables

- Can introduce subtle and confusing bugs!
- Limit use of global variables to situations in which they're really needed.
  - Shared variables.



$a_{00}$	$a_{01}$	$\cdots$	$a_{0,n-1}$
$a_{10}$	$a_{11}$	$\cdots$	$a_{1,n-1}$
$\vdots$	$\vdots$		$\vdots$
$a_{i0}$	$a_{i1}$	$\cdots$	$a_{i,n-1}$
$\vdots$	$\vdots$		$\vdots$
$a_{m-1,0}$	$a_{m-1,1}$	$\cdots$	$a_{m-1,n-1}$

$\begin{matrix} x_0 \\ x_1 \\ \vdots \\ x_{n-1} \end{matrix} =$

$y_0$
$y_1$
$\vdots$
$y_i = a_{i0}x_0 + a_{i1}x_1 + \cdots a_{i,n-1}x_{n-1}$
$\vdots$
$y_{m-1}$

# MATRIX-VECTOR MULTIPLICATION IN PTHREADS

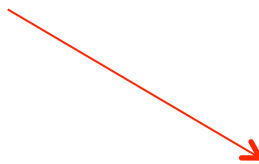
# Serial pseudo-code

```
/* For each row of A */  
for (i = 0; i < m; i++) {  
    y[i] = 0.0;  
    /* For each element of the row and each element of x */  
    for (j = 0; j < n; j++)  
        y[i] += A[i][j]* x[j];  
}
```

$$y_i = \sum_{j=0}^{n-1} a_{ij}x_j$$

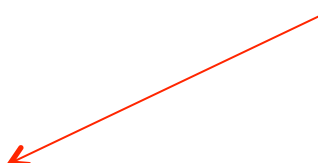
# Using Pthreads

Thread	Components of $y$
0	$y[0], y[1]$
1	$y[2], y[3]$
2	$y[4], y[5]$



```
y[0] = 0.0;
for (j = 0; j < n; j++)
    y[0] += A[0][j]* x[j];
```

thread 0



```
y[i] = 0.0;
for (j = 0; j < n; j++)
    y[i] += A[i][j]* x[j];
```

Thread  $i$



# Pthreads matrix-vector multiplication

```
void *nth_mat_vect(void *rank) {
    int *rank_int_ptr = (int*) rank;
    int my_rank = *rank_int_ptr;
    int local_m = M / NUM_THREADS;
    int my_first_row = (my_rank * local_m);
    int my_last_row = (my_rank+1)*local_m - 1;

    for (int i = my_first_row; i <= my_last_row; i++) {
        Y[i] = 0.0;
        for (int j = 0; j < N; j++) {
            Y[i] += A[i][j] * X[j];
        }
    }
    return NULL;
}
```



# CRITICAL SECTIONS

# Estimating $\pi$

$$\pi = 4 \left( 1 - \frac{1}{3} + \frac{1}{5} - \frac{1}{7} + \cdots + (-1)^n \frac{1}{2n+1} + \cdots \right)$$

```
double factor = 1.0;  
double sum = 0.0;  
for (i = 0; i < n; i++, factor = -factor) {  
    sum += factor/(2*i+1);  
}  
pi = 4.0*sum;
```

# A thread function for computing $\pi$

```
void *calc(void *rank) {
    int *rank_int_ptr = (int*) rank;
    int my_rank = *rank_int_ptr;
    int my_n = N / NUM_THREADS;
    int my_first_i = (my_rank * my_n);
    int my_last_i = my_first_i + my_n;
    double factor = -1.0;
    if (my_first_i % 2 == 0) {
        factor = 1.0;
    }
    for (int i = my_first_i; i <= my_last_i; i++) {
        sum += factor / (2*i+1);
        factor = -1 * factor;
    }
    return NULL;
}
```

# Using a dual core processor

	$n$			
	$10^5$	$10^6$	$10^7$	$10^8$
$\pi$	3.14159	3.141593	3.1415927	3.14159265
1 Thread	3.14158	3.141592	3.1415926	3.14159264
2 Threads	3.14158	3.141480	3.1413692	3.14164686

Note that as we increase  $n$ , the estimate with one thread gets better and better.

Why?



# Possible race condition

```
int x = 0;
void *compute(void *rank) {
    int *rank_int_ptr = (int*) rank;
    int y = *rank_int_ptr;
    x += y;
}
```

Time	Thread 0	Thread 1
1	Started by main thread	
2	Call Compute ()	Started by main thread
3	Assign y = 1	Call Compute ()
4	Put x=0 and y=1 into registers	Assign y = 2
5	Add 0 and 1	Put x=0 and y=2 into registers
6	Store 1 in memory location x	Add 0 and 2
7		Store 2 in memory location x

# Busy-Waiting

- A thread repeatedly tests a condition, but, effectively, does no useful work until the condition has the appropriate value.
- Beware of optimizing compilers, though!

```
y = Compute(my_rank);  
while (flag != my_rank);  
x = x + y;  
flag++;
```

flag initialized to 0 by main thread

# Pthreads global sum with busy-waiting

```
void* Thread_sum(void* rank) {
    long my_rank = (long) rank;
    double factor;
    long long i;
    long long my_n = n/thread_count;
    long long my_first_i = my_n*my_rank;
    long long my_last_i = my_first_i + my_n;

    if (my_first_i % 2 == 0)
        factor = 1.0;
    else
        factor = -1.0;

    for (i = my_first_i; i < my_last_i; i++, factor = -factor) {
        while (flag != my_rank);
        sum += factor/(2*i+1);
        flag = (flag+1) % thread_count;
    }

    return NULL;
} /* Thread_sum */
```



# Busy-waiting sum

```
void *calc(void *rank) {
    int *rank_int_ptr = (int*) rank;
    int my_rank = *rank_int_ptr;
    int my_n = N / NUM_THREADS;
    int my_first_i = (my_rank * my_n);
    int my_last_i = my_first_i + my_n;
    double factor = -1.0;
    if (my_first_i % 2 == 0) {
        factor = 1.0;
    }
    for (int i = my_first_i; i <= my_last_i; i++) {
        while (flag != my_rank);
        sum += factor / (2*i+1);
        flag = (flag+1) % NUM_THREADS;
        factor = -1 * factor;
    }
    return NULL;
}
```

# Time?

With no barriers – 2 threads  
pi=3.134463

real 0m0.739s  
**user 0m1.445s**  
sys 0m0.012s

With busy-wait – 2 threads  
pi=3.141593

real 0m4.509s  
**user 0m8.815s**  
sys 0m0.053s

With no barriers – 3 threads  
pi=3.146096

real 0m0.647s  
**user 0m1.792s**  
sys 0m0.006s

With no barriers – 3 threads  
pi=3.141593

real 0m6.872s  
**user 0m19.383s**  
sys 0m0.056s

# Busy-waiting sum at end

```
void *calc(void *rank) {
    int *rank_int_ptr = (int*) rank;
    int my_rank = *rank_int_ptr;
    int my_n = N / NUM_THREADS;
    int my_first_i = (my_rank * my_n);
    int my_last_i = my_first_i + my_n;
    double factor = -1.0;
    if (my_first_i % 2 == 0) {
        factor = 1.0;
    }
    double my_sum = 0.0;
    for (int i = my_first_i; i <= my_last_i; i++) {
        my_sum += factor / (2*i+1);
        factor = -1 * factor;
    }
    while (flag != my_rank);
    sum += my_sum;
    flag = (flag+1) % NUM_THREADS;
    return NULL;
}
```

# Mutexes

- A thread that is busy-waiting may continually use the CPU accomplishing nothing.
- Mutex (mutual exclusion) is a special type of variable that can be used to restrict access to a critical section to a single thread at a time.

# Mutexes



- Used to guarantee that one thread “excludes” all other threads while it executes the critical section.
- The Pthreads standard includes a special type for mutexes: `pthread_mutex_t`.

# Pthread Mutexes

- Type: `pthread_mutex_t`

```
int pthread_mutex_init(pthread_mutex_t *mutex,  
                        const pthread_mutexattr_t *attr);  
int pthread_mutex_destroy(pthread_mutex_t *mutex);  
int pthread_mutex_lock(pthread_mutex_t *mutex);  
int pthread_mutex_unlock(pthread_mutex_t *mutex);  
int pthread_mutex_trylock(pthread_mutex_t *mutex);
```

- Important: Mutex scope must be visible to all threads!

# Busy-waiting sum at end

```
void *calc(void *rank) {
    int *rank_int_ptr = (int*) rank;
    int my_rank = *rank_int_ptr;
    int my_n = N / NUM_THREADS;
    int my_first_i = (my_rank * my_n);
    int my_last_i = my_first_i + my_n;
    double factor = -1.0;
    if (my_first_i % 2 == 0) {
        factor = 1.0;
    }
    double my_sum = 0.0;
    for (int i = my_first_i; i <= my_last_i; i++) {
        my_sum += factor / (2*i+1);
        factor = -1 * factor;
    }
    pthread_mutex_lock(&mutex);
    sum += my_sum;
    pthread_mutex_unlock(&mutex);
    return NULL;
}
```

Threads	Busy-Wait	Mutex
1	2.90	2.90
2	1.45	1.45
4	0.73	0.73
8	0.38	0.38
16	0.50	0.38
32	0.80	0.40
64	3.56	0.38

$$\frac{T_{\text{serial}}}{T_{\text{parallel}}} \approx \text{thread\_count}$$

Run-times (in seconds) of  $\pi$  programs using  $n = 108$  terms on a system with two four-core processors.



Time	flag	Thread				
		0	1	2	3	4
0	0	crit sect	busy wait	susp	susp	susp
1	1	terminate	crit sect	susp	busy wait	susp
2	2	—	terminate	susp	busy wait	busy wait
⋮	⋮			⋮	⋮	⋮
?	2	—	—	crit sect	susp	busy wait

Possible sequence of events with busy-waiting and more threads than cores.

# Issues

- Busy-waiting enforces the order threads access a critical section.
- Using mutexes, the order is left to chance and the system.
- There are applications where we need to control the order threads access the critical section.
  - Producer/Consumer (message sending)

# Semaphores

- Similar to mutexes but there isn't necessarily ownership of "data", but coordination.
- Think of it as a counter:
  - When counter = 0 its locked (wait)
  - When counter > 0 its unlocked (proceed)
- Two main functions:
  - Wait – decrements the counter if > 0 returns; if  $\leq 0$  then waits
  - Post – increments the counter
- Not provided natively using pthreads, although fairly straight forward to implement.

# Message Sending Example

```
Void * send_msg(void* rank) {
    int *rank_int_ptr = (int*) rank;
    int my_rank = *rank_int_ptr;
    int dest = my_rank + 1 % NUM_THREADS;
    char* my_msg = malloc(MSG_MAX*sizeof(char));
    sprintf(my_msg, "Hello from %d", my_rank);

    messages[dest] = my_msg;
    sem_post(&semaphores[dest]);

    sem_wait(&semaphores[my_rank])
    printf("Thread %d> %s\n", my_rank, messages[my_rank]);

    return NULL;
}
```

# Barriers

- Synchronizing the threads to make sure that they all are at the same point in a program is called a barrier.
- No thread can cross the barrier until all the threads have reached it.

# Busy-waiting and a Mutex

- Implementing a barrier using busy-waiting and a mutex is straightforward.
- We use a shared counter protected by the mutex.
- When the counter indicates that every thread has entered the critical section, threads can leave the critical section.

# Why Barriers

- Synchronization at algorithm steps.
  - Iterations in linear solver, check to see CI is good.
- Write state to disk
  - Simulated wind tunnel, write image.
  - Checkpointing
- Debugging
  - Print state of matrixes at different steps
- Error Checking
  - Verify algorithm is converging

# Barriers with Mutex

```
void* Thread_work( . . . ) {  
    . . .  
  
    /* START Barrier */  
    pthread_mutex_lock(&barrier_mutex);  
    counter++;  
    pthread_mutex_unlock(&barrier_mutex);  
    while (counter < thread_count);  
    /* END Barrier */  
  
    . . .  
}
```



# Condition Variables

- A condition variable is a data object that allows a thread to suspend execution until a certain event or condition occurs.
- When the event or condition occurs another thread can signal the thread to “wake up.”
- A condition variable is always associated with a mutex.

# Condition variables

- Type `pthread_cond_t`

```
int pthread_cond_init(pthread_cond_t *cond,  
                      const pthread_condattr_t *attr);  
int pthread_cond_destroy(pthread_cond_t *cond);  
int pthread_cond_wait(pthread_cond_t *cond,  
                      pthread_mutex_t *mutex);  
int pthread_cond_signal(pthread_cond_t *cond);  
int pthread_cond_broadcast(pthread_cond_t *cond);
```

# Barrier With Condition

```
void* Thread_work( . . . ) {
    . . .

    /* START Barrier */
    pthread_mutex_lock(&mutex);
    counter++;
    if (counter == thread_count) {
        counter = 0;
        pthread_cond_broadcast(&cond_var);
    } else {
        while(pthread_cond_wait(&cond_var, &mutex) != 0);
    }
    pthread_mutex_unlock(&mutex);
    /* END Barrier */

    . . .
}
```

# Some rules...

- Shared data should always be accessed through a single mutex
- Think of a boolean condition (expressed in terms of program variables) for each condition variable. Every time the value of the boolean condition may have changed, call `Broadcast` for the condition variable
  - Only call `Signal` when you are absolutely certain that *any and only one* waiting thread can enter the critical section
  - if noone is waiting, signal is lost
- Signaling/Broadcasting thread need not have the mutex
  - may be more efficient to release it first...
- Globally order locks, acquire in order in all threads




# **READ-WRITE LOCKS**

## **SECTION 4.9**

# Controlling access to a large, shared data structure

- Let's look at an example.
- Suppose the shared data structure is a sorted linked list of ints, and the operations of interest are Member, Insert, and Delete.

# Linked Lists




```
/* Struct for list nodes */
struct list_node_s {
    int      data;
    struct list_node_s* next;
};
```

# Linked List Membership

```
int  Member(int value) {  
    struct list_node_s* temp;  
  
    temp = head;  
    while (temp != NULL && temp->data < value)  
        temp = temp->next;  
  
    if (temp == NULL || temp->data > value) {  
        return 0;  
    } else {  
        return 1;  
    }  
}
```



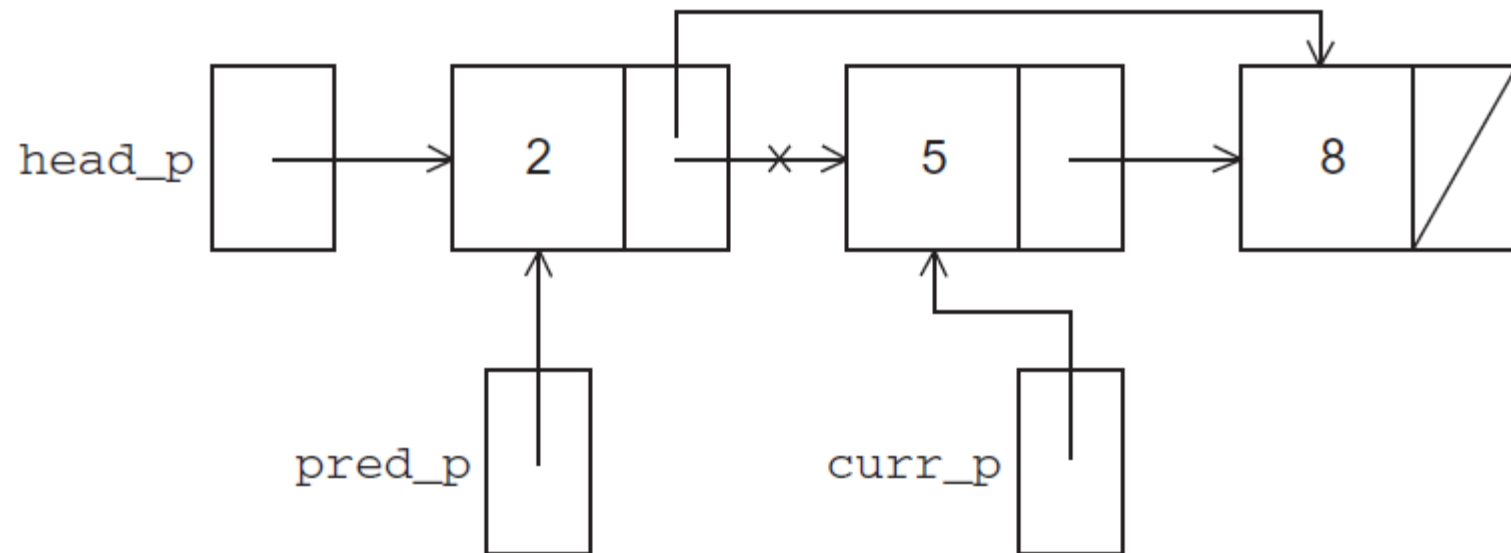
# Inserting a new node into a list



# Inserting a new node into a list

```
int Insert(int value) {
    struct list_node_s* curr = head;
    struct list_node_s* pred = NULL;
    struct list_node_s* temp;
    while (curr != NULL && curr->data < value) {
        pred = curr;
        curr = curr->next;
    }
    if (curr == NULL || curr->data > value) {
        temp = malloc(sizeof(struct list_node_s));
        temp->data = value;
        temp->next = curr;
        if (pred == NULL)
            head = temp;
        else
            pred->next = temp;
    } else { /* value in list */
        return 0;
    }
}
```

# Deleting a node from a linked list



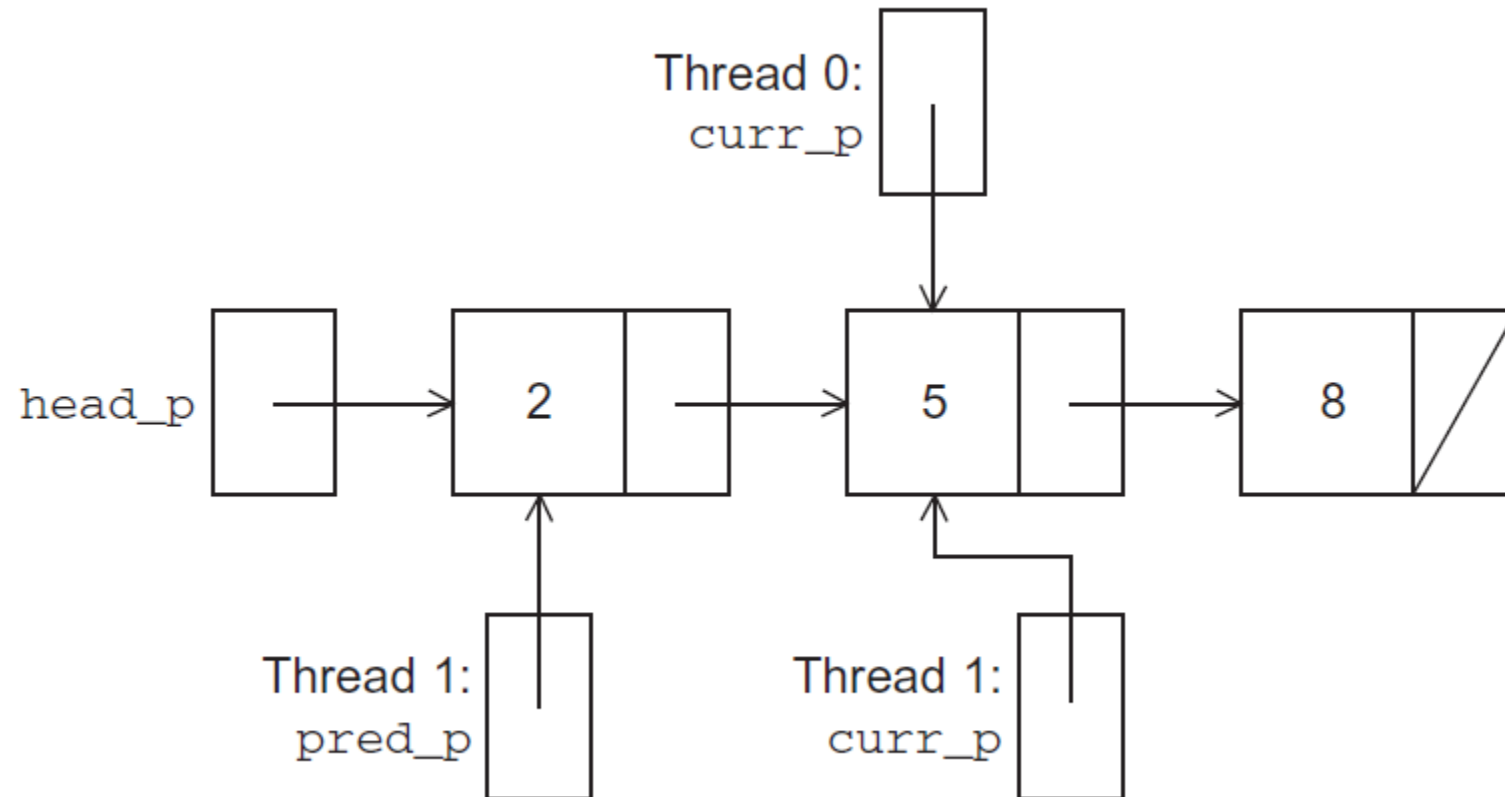
# Deleting a node from a linked list

```
int Delete(int value) {
    struct list_node_s* curr = head;
    struct list_node_s* pred = NULL;
    int rv = 1;
    /* Find value */
    while (curr != NULL && curr->data < value) {
        pred = curr;
        curr = curr->next;
    }
    if (curr != NULL && curr->data == value) {
        if (pred == NULL) { /* first element in list */
            head = curr->next;
            free(curr);
        } else {
            pred->next = curr->next;
            free(curr);
        }
        return 1;
    } else { /* Not in list */
        return 0;
    }
}
```

# A Multi-Threaded Linked List

- Let's try to use these functions in a Pthreads program.
- In order to share access to the list, we can define `head_p` to be a global variable.
- This will simplify the function headers for `Member`, `Insert`, and `Delete`, since we won't need to pass in either `head_p` or a pointer to `head_p`: we'll only need to pass in the value of interest.

# Simultaneous access by two threads



# Solution #1

- An obvious solution is to simply lock the list any time that a thread attempts to access it.
- A call to each of the three functions can be protected by a mutex.

```
Pthread_mutex_lock(&list_mutex);  
Member(value);  
Pthread_mutex_unlock(&list_mutex);
```

# Issues

- We're serializing access to the list.
- If the vast majority of our operations are calls to **Member**, we'll fail to exploit this opportunity for parallelism.
- On the other hand, if most of our operations are calls to **Insert** and **Delete**, then this may be the best solution since we'll need to serialize access to the list for most of the operations, and this solution will certainly be easy to implement.



## Solution #2

- Instead of locking the entire list, we could try to lock individual nodes.
- A “finer-grained” approach.

```
/* Struct for list nodes */
struct list_node_s {
    int      data;
    struct list_node_s* next;
    pthread_mutex_t mutex;
};
```

# Issues

- This is much more complex than the original **Member** function.
- It is also much slower, since, in general, each time a node is accessed, a mutex must be locked and unlocked.
- The addition of a mutex field to each node will substantially increase the amount of storage needed for the list.

## Implementation of Member with one mutex per list node (1)

```
int Member(int value) {  
    struct list_node_s* temp_p;  
  
    pthread_mutex_lock(&head_p_mutex);  
    temp_p = head_p;  
    while (temp_p != NULL && temp_p->data < value) {  
        if (temp_p->next != NULL)  
            pthread_mutex_lock(&(temp_p->next->mutex));  
        if (temp_p == head_p)  
            pthread_mutex_unlock(&head_p_mutex);  
        pthread_mutex_unlock(&(temp_p->mutex));  
        temp_p = temp_p->next;  
    }  
}
```

## Implementation of Member with one mutex per list node (2)

```
    if (temp_p == NULL || temp_p->data > value) {
        if (temp_p == head_p)
            pthread_mutex_unlock(&head_p_mutex);
        if (temp_p != NULL)
            pthread_mutex_unlock(&(temp_p->mutex));
        return 0;
    } else {
        if (temp_p == head_p)
            pthread_mutex_unlock(&head_p_mutex);
        pthread_mutex_unlock(&(temp_p->mutex));
        return 1;
    }
} /* Member */
```

# Pthreads Read-Write Locks

- Neither of our multi-threaded linked lists exploits the potential for simultaneous access to any node by threads that are executing Member.
- The first solution only allows one thread to access the entire list at any instant.
- The second only allows one thread to access any given node at any instant.

# Pthreads Read-Write Locks

- A read-write lock is somewhat like a mutex except that it provides two lock functions.
- The first lock function locks the read-write lock for reading, while the second locks it for writing.
- So multiple threads can simultaneously obtain the lock by calling the read-lock function, while only one thread can obtain the lock by calling the write-lock function.
- Thus, if any threads own the lock for reading, any threads that want to obtain the lock for writing will block in the call to the write-lock function.
- If any thread owns the lock for writing, any threads that want to obtain the lock for reading or writing will block in their respective locking functions.

# Protecting our linked list

```
pthread_rwlock_rdlock(&rwlock);  
Member(val);  
pthread_rwlock_unlock(&rwlock);
```

```
pthread_rwlock_wrlock(&rwlock);  
Insert(val);  
pthread_rwlock_unlock(&rwlock);
```

```
pthread_rwlock_wrlock(&rwlock);  
Delete(val);  
pthread_rwlock_unlock(&rwlock);
```

# Linked List Performance

Implementation	Number of Threads			
	1	2	4	8
Read-Write Locks	0.213	0.123	0.098	0.115
One Mutex for Entire List	0.211	0.450	0.385	0.457
One Mutex per Node	1.680	5.700	3.450	2.700

100,000 ops/thread

99.9% Member

0.05% Insert

0.05% Delete



# Linked List Performance

Implementation	Number of Threads			
	1	2	4	8
Read-Write Locks	2.48	4.97	4.69	4.71
One Mutex for Entire List	2.50	5.13	5.04	5.11
One Mutex per Node	12.00	29.60	17.00	12.00

100,000 ops/thread

80% Member

10% Insert

10% Delete



# **THREAD-SAFETY**

## **SECTION 4.11**


# Thread-Safety

- A block of code is **thread-safe** if it can be simultaneously executed by multiple threads without causing problems



# Example

- Suppose we want to use multiple threads to “tokenize” a file that consists of ordinary English text.
- The tokens are just contiguous sequences of characters separated from the rest of the text by white-space — a space, a tab, or a newline.



# Simple approach

- Divide the input file into lines of text and assign the lines to the threads in a round-robin fashion.
- The first line goes to thread 0, the second goes to thread 1, . . . , the  $t$ th goes to thread  $t$ , the  $t + 1$ st goes to thread 0, etc.

# Simple approach

- We can serialize access to the lines of input using mutexes.
- After a thread has read a single line of input, it can tokenize the line using the `strtok` function.

# The strtok function

- The first time it's called the string argument should be the text to be tokenized.
  - Our line of input.
- For subsequent calls, the first argument should be NULL.

```
char * strtok (char* string,  
               const char* separators)
```

# The strtok function

- The idea is that in the first call, `strtok` caches a pointer to string, and for subsequent calls it returns successive tokens taken from the cached copy.



# Multi-threaded tokenizer

```
void *Tokenize(void* rank) {
    . . .
    /* Force sequential reading of the input */
    pthread_mutex_lock(&mutex);
    fg_rv = fgets(my_line, MAX, fp);
    pthread_mutex_unlock(&mutex);
    while (fg_rv != NULL) {
        printf("Thread %d > my line = %s", my_rank, my_line);
        count = 0;
        my_string = strtok(my_line, " \t\n");
        while ( my_string != NULL ) {
            count++;
            printf("Thread %d > string %d = %s\n", my_rank, count, my_string);
            my_string = strtok(NULL, " \t\n");
        }
        printf("Thread %d > After tokenizing, my_line = %s\n", my_rank, my_line);
        pthread_mutex_lock(&mutex);
        fg_rv = fgets(my_line, MAX, fp);
        pthread_mutex_unlock(&mutex);
    }
    . . .
}
```

# What happened?

- `strtok` caches the input line by declaring a variable to have static storage class.
- This causes the value stored in this variable to persist from one call to the next.
- Unfortunately for us, this cached string is shared, not private.

# What happened?

- Thus, thread 0's call to `strtok` with the third line of the input has apparently overwritten the contents of thread 1's call with the second line.
- So the `strtok` function is not thread-safe. If multiple threads call it simultaneously, the output may not be correct.

# Other unsafe C library functions

- Regrettably, it's not uncommon for C library functions to fail to be thread-safe.
- The random number generator `random` in `stdlib.h`.
- The time conversion function `localtime` in `time.h`.

# “re-entrant” functions

- In some cases, the C standard specifies an alternate, thread-safe, version of a function.

```
char * strtok (char* string,  
               const char* separators,  
               char** saveptr_p)
```

- Not all re-entrant functions are threadsafe!

# Python and GIL

- Python using a **global interpreter lock**
- This means that any python code can't truly be executed in parallel
  - Although you can use threads in the c code and explicitly unlock other processes
  - Basic I/O functions do this by default, therefore multiple python threads doing I/O do not block one another (ie webserver, file reads)
- Other implementations of python do not have this constraint (JPython)
- Ruby has some similar constraints

# Concluding Remarks (1)

- A thread in shared-memory programming is analogous to a process in distributed memory programming.
- However, a thread is often lighter-weight than a full-fledged process.
- In Pthreads programs, all the threads have access to global variables, while local variables usually are private to the thread running the function.

# Concluding Remarks (2)

- When indeterminacy results from multiple threads attempting to access a shared resource such as a shared variable or a shared file, at least one of the accesses is an update, and the accesses can result in an error, we have a **race condition**.



# Concluding Remarks (3)

- A **critical section** is a block of code that updates a shared resource that can only be updated by one thread at a time.
- So the execution of code in a critical section should, effectively, be executed as serial code.

# Concluding Remarks (4)

- **Busy-waiting** can be used to avoid conflicting access to critical sections with a flag variable and a while-loop with an empty body.
- It can be very wasteful of CPU cycles.
- It can also be unreliable if compiler optimization is turned on.

# Concluding Remarks (5)

- A **mutex** can be used to avoid conflicting access to critical sections as well.
- Think of it as a lock on a critical section, since mutexes arrange for mutually exclusive access to a critical section.

# Concluding Remarks (6)

- A **semaphore** is the third way to avoid conflicting access to critical sections.
- It is an unsigned int together with two operations: `sem_wait` and `sem_post`.
- Semaphores are more powerful than mutexes since they can be initialized to any nonnegative value.

# Concluding Remarks (7)

- A **barrier** is a point in a program at which the threads block until all of the threads have reached it.
- A **read-write lock** is used when it's safe for multiple threads to simultaneously read a data structure, but if a thread needs to modify or write to the data structure, then only that thread can access the data structure during the modification.

# Concluding Remarks (8)

- Some C functions cache data between calls by declaring variables to be static, causing errors when multiple threads call the function.
- This type of function is not **thread-safe**.