

Parallel Algorithms & the PRAM Model

Advanced Topics Spring 2009

Prof. Robert van Engelen






Overview

- The PRAM model of parallel computation
- Simulations between PRAM models
- Work-time presentation framework of parallel algorithms
- Design and analysis of parallel algorithms



The PRAM Model of Parallel Computation

- Parallel Random Access Machine (PRAM)
- Natural extension of RAM: each processor is a RAM
- Processors operate synchronously
- Earliest and best-known model of parallel computation



Shared memory with m locations

p processors, each with private memory

All processors operate synchronously, by executing load, store, and operations on data



Synchronous PRAM versus Asynchronous PRAM

- The **synchronous PRAM** model has a similarity with data-parallel execution on a SIMD machine
 - All processors execute the same program
 - All processors execute the same PRAM step instruction stream in “lock-step”
 - Effect of operation depends on local data
 - Instructions can be selectively disabled (for if-then-else flow)

- The **asynchronous PRAM** model
 - Several competing models
 - No lock-step



Classification of PRAM Model

- A PRAM step (“clock cycle”) consists of three phases
 1. *Read*: each processor may read a value from shared memory
 2. *Compute*: each processor may perform operations on local data
 3. *Write*: each processor may write a value to shared memory
- Model is refined for concurrent read/write capability
 - Exclusive Read Exclusive Write (EREW)
 - Concurrent Read Exclusive Write (CREW)
 - Concurrent Read Concurrent Write (CRCW)
- CRCW PRAM: what to do with concurrent writes?
 - Common CRCW: all processors must write the same value
 - Arbitrary CRCW: one of the processors succeeds in writing
 - Priority CRCW: processor with highest priority succeeds in writing



Comparison of PRAM Models

- A model A is less powerful compared to model B if either
 - The time complexity is asymptotically less in model B for solving a problem compared to A
 - Or the time complexity is the same and the work complexity is asymptotically less in model B compared to A
- From weakest to strongest:
 - EREW
 - CREW
 - Common CRCW
 - Arbitrary CRCW
 - Priority CRCW



Simulations Between PRAM Models

- An algorithm designed for a weaker model can be executed within the same time complexity and work complexity on a stronger model

- An algorithm designed for a stronger model can be *simulated* on a weaker model, either with
 - Asymptotically more processors (or more work by the same number of processors)
 - Or asymptotically more time



Simulating a Priority CRCW on an EREW PRAM


- Theorem: An algorithm that runs in T time on the p -processor priority CRCW PRAM can be simulated by EREW PRAM to run in $O(T \log p)$ time
 - A concurrent read or write of an p -processor CRCW PRAM can be implemented on a p -processor EREW PRAM to execute in $O(\log p)$ time
 - Q_1, \dots, Q_p CRCW processors, such that Q_i has to read (write) $M[j_i]$
 - P_1, \dots, P_p EREW processors
 - M_1, \dots, M_p denote shared memory locations for special use
 - P_i stores $\langle j_i, i \rangle$ in M_i
 - Sort pairs in lexicographically non-decreasing order in $O(\log p)$ time using EREW merge sort algorithm
 - Pick representative from each block of pairs that have same first component in $O(1)$ time
 - Representative P_i reads (writes) from $M[k]$ with $\langle k, _ \rangle$ in M_i and copies data to each M in the block in $O(\log p)$ time using EREW segmented parallel prefix algorithm
 - P_i reads data from M_i



Example 1:

Reduction on the EREW PRAM

- Reduce (sum) p values on the p -processor EREW PRAM in $O(\log p)$ time
- Reduction algorithm uses exclusive reads and writes
- Algorithm is the basis of other EREW algorithms





Example 1

Sum of n values using n processors (i)

Each processor i , $1 \leq i \leq n$, executes:

Input: $A[1, \dots, n]$, $n = 2^k$

Output: sum $S = \sum_{j=1..n} A[j]$

begin

$B[i] := A[i]$

for $h = 1$ **to** $\log n$ **do**

if $i \leq n/2^h$ **then**

$B[i] := B[2i-1] + B[2i]$

if $i = 1$ **then**

$S := B[i]$

end

How much time?

How many operations?



Example 2: Matrix Multiply on the CREW PRAM

- Consider $n \times n$ matrix multiplication $C = A B$ using n^3 processors
- Each element of C

$$c_{ij} = \sum_{k=1..n} a_{ik} b_{kj}$$

can be computed on the CREW PRAM in parallel using n processors in $O(\log n)$ time

- All c_{ij} can be computed using n^3 processors in $O(\log n)$ time



Example 2

Matrix multiply with n^3 processors (i,j,l)

Each processor (i,j,l) executes:

Input: $n \times n$ matrices A and B , $n = 2^k$

Output: $C = A B$

begin

$C'[i,j,l] := A[i,l]B[l,j]$

for $h = 1$ **to** $\log n$ **do**

if $i \leq n/2^h$ **then**

$C'[i,j,l] := C'[i,j,2l-1] + C'[i,j,2l]$

if $l = 1$ **then**

$C[i,j] := C'[i,j,1]$

end

$O(\log n)$ time

How many operations?



Example 2: CREW versus EREW PRAM

- Algorithm on the CREW PRAM requires $O(\log n)$ time and $O(n^3)$ operations (n^2 processors perform $O(n)$ ops)
- On the EREW PRAM, the exclusive reads of a_{ij} and b_{ij} values can be satisfied by making n copies of a and b , which takes $O(\log n)$ time with n processors (broadcast tree)
- Total time is still $O(\log n)$
- But requires more work and total memory requirement is huge!



The WT Scheduling Principle

- The **work-time (WT) scheduling principle** schedules p processors to execute an algorithm
 - Algorithm has $T(n)$ time steps and $W(n)$ total operations
 - A time step can be parallel, i.e. **pardo**
- We can adapt the algorithm to run on the p -processor PRAM in $\leq \lfloor W(n)/p \rfloor + T(n)$ steps
- Proof
 - Let $W_i(n)$ be the number of operations (work) performed in time unit i , $1 \leq i \leq T(n)$
 - Simulate each set of $W_i(n)$ operations in $\lceil W_i(n)/p \rceil$ parallel steps, for each $1 \leq i \leq T(n)$
 - The number of steps on the p -processor PRAM takes
$$\sum_i \lceil W_i(n)/p \rceil \leq \sum_i (\lfloor W_i(n)/p \rfloor + 1) \leq \lfloor W(n)/p \rfloor + T(n)$$



Work-Time Presentation

- The WT presentation can be used to determine the time and operation requirements of an algorithm
- The upper-level WT presentation framework describes the algorithm in terms of a sequence of time units
 - From which we can determine $T(n)$ and $W(n)$
- The lower-level follows the WT scheduling principle
 - p -processor PRAM requires $\leq \lfloor W(n)/p \rfloor + T(n)$ steps



Example 1 Revisited: WT Presentation

```
Input:  $A[1, \dots, n]$ ,  $n = 2^k$   
Output:  $\text{sum } S = \sum_{j=1..n} A[j]$   
begin  
  for  $1 \leq i \leq n$  pardo  
     $B[i] := A[i]$   
  for  $h = 1$  to  $\log n$  do  
    for  $1 \leq i \leq n/2^h$  pardo  
       $B[i] := B[2i-1] + B[2i]$   
    if  $i = 1$  then  
       $S := B[1]$   
end
```

*Do you spot any
concurrent reads?
concurrent writes?*

$$T(n) = O(\log n)$$
$$W(n) = O(n)$$

WT scheduling principle:
total time $\leq O(n/p + \log n)$



Example 2 Revisited: WT-Presentation

```
Input:  $n \times n$  matrices  $A$  and  $B$ ,  $n = 2^k$   
Output:  $C = A B$   
begin  
  for  $1 \leq i, j, l \leq n$  pardo  
     $C'[i,j,l] := A[i,l]B[l,j]$   
  for  $h = 1$  to  $\log n$  do  
    for  $1 \leq i, j \leq n, 1 \leq l \leq n/2^h$  pardo  
       $C'[i,j,l] := C'[i,j,2l-1] + C'[i,j,2l]$   
    for  $1 \leq i, j \leq n$  pardo  
       $C[i,j] := C'[i,j,1]$   
end
```

$$T(n) = O(\log n)$$

$$W(n) = n^3$$

WT scheduling principle:
total time $\leq O(n^3/p + \log n)$



Example 3: PRAM Recursive Prefix Sum Algorithm

Input: Array of (x_1, x_2, \dots, x_n) elements, $n = 2^k$

Output: Prefix sums s_i , $1 \leq i \leq n$

begin

if $n = 1$ **then** $s_1 = x_1$; **exit**

for $1 \leq i \leq n/2$ **pardo**

$y_i := x_{2i-1} + x_{2i}$

 Recursively compute prefix sums of y and store in z

for $1 \leq i \leq n$ **pardo**

if i is even **then** $s_i := z_{i/2}$

if $i > 1$ is odd **then** $s_i := z_{(i-1)/2} + x_i$

if $i = 1$ **then** $s_1 := x_1$

end



Proof of Work Optimality

- **Theorem:** The PRAM prefix sum algorithm correctly computes the prefix sum and takes $T(n) = O(\log n)$ time using a total of $W(n) = O(n)$ operations
- **Proof** by induction on k , where input size $n = 2^k$
 - Base case $k = 0$: $s_1 = x_1$
 - Assume correct for $n = 2^k$
 - For $n = 2^{k+1}$
 - For all $1 \leq j \leq n/2$ we have
$$z_j = y_1 + y_2 + \dots + y_j = (x_1 + x_2) + (x_3 + x_4) + \dots + (x_{2j-1} + x_{2j})$$
 - Hence, for $i = 2j \leq n$ we have $s_i = s_{2j} = z_j = z_{i/2}$
 - And $i = 2j+1 \leq n$ we have $s_i = s_{2j+1} = s_{2j} + x_{2j+1} = z_j + x_{2j+1} = z_{(i-1)/2} + x_i$
- $T(n) = T(n/2) + a \quad \Rightarrow T(n) = O(\log n)$
- $W(n) = W(n/2) + bn \quad \Rightarrow W(n) = O(n)$




PRAM Nonrecursive Prefix Sum

```
Input: Array  $A$  of size  $n = 2^k$   
Output: Prefix sums in  $C[0,j]$ ,  $1 \leq j \leq n$   
begin  
  for  $1 \leq j \leq n$  pardo  
     $B[0,j] := A[j]$   
  for  $h = 1$  to  $\log n$  do  
    for  $1 \leq j \leq n/2^h$  pardo  
       $B[h,j] := B[h-1,2j-1] + B[h-1,2j]$   
  for  $h = \log n$  to  $0$  do  
    for  $1 \leq j \leq n/2^h$  pardo  
      if  $j$  is even then  $C[h,j] := C[h+1,j/2]$   
      else if  $i = 1$  then  $C[h,1] := B[h,1]$   
      else  $C[h,j] := C[h+1,(j-1)/2] + B[h,j]$   
end
```

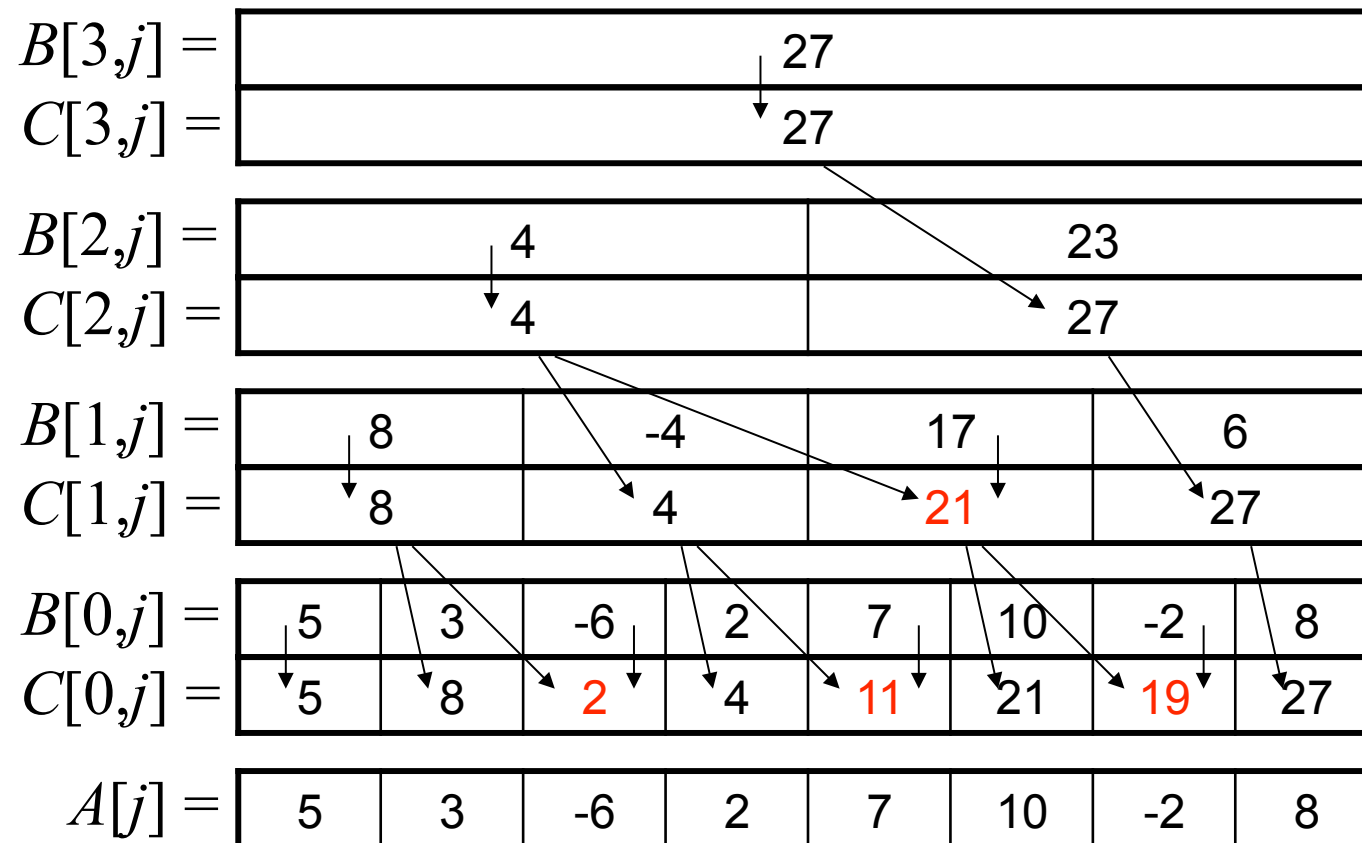


First Pass: Bottom-Up






Second Pass: Top-Down





Example 4: Pointer Jumping

- Finding the roots of a forest using pointer-jumping





Pointer Jumping on the CREW PRAM

Input: A forest of trees, each with a self-loop at its root, consisting of arcs $(i, P(i))$ and nodes i , where $1 \leq i \leq n$

Output: For each node i , the root $S[i]$

begin

for $1 \leq i \leq n$ **pardo**

$S[i] := P[i]$

while $S[i] \neq S[S[i]]$ **do**

$S[i] := S[S[i]]$

end

$T(n) = O(\log h)$ with h the maximum height of trees

$W(n) = O(n \log h)$



PRAM Model Summary

- PRAM removes algorithmic details concerning synchronization and communication, allowing the algorithm designer to focus on problem properties
- A PRAM algorithm includes an explicit understanding of the operations performed at each time unit and an explicit allocation of processors to jobs at each time unit
- PRAM design paradigms have turned out to be robust and have been mapped efficiently onto many other parallel models and even network models
 - A SIMD network model considers *communication diameter*, *bisection width*, and *scalability* properties of the network topology of a parallel machine such as a mesh or hypercube



Design and Analysis of Parallel Algorithms

- Arithmetic problems:
 - Polynomial evaluation: first-order linear recurrence
 - Polynomial multiplication: FFT
 - Lagrange interpolation

- Planar geometry:
 - The convex hull problem revisited: constant-time computation of the upper common tangent



First-Order Linear Recurrences

- Consider the first-order linear recurrence:

$$y_1 = b_1$$

$$y_i = a_i y_{i-1} + b_i \quad \text{for } 2 \leq i \leq n$$

- At first sight this seems impossible to parallelize, at least in its current form
- However, note that the prefix sum

$$y_i = \sum_{j=1..i} b_j$$

is a special case of a first-order linear recurrence where $a_i = 1$ (the multiplicative unit element)

- We know how to parallelize the prefix sum



Divide and Conquer Parallelization

- Rewrite $y_i = a_i y_{i-1} + b_i$ into $y_i = a_i (a_{i-1} y_{i-2} + b_{i-1}) + b_i$
- This equation defines a linear recurrence of size $n/2$ for even index i

$$\begin{aligned} z_1 &= b_1' \\ z_i &= a_i' z_{i-1} + b_i' \quad 2 \leq i \leq n/2 \end{aligned}$$

1. Let

$$\begin{aligned} a_i' &= a_{2i} a_{2i-1} \\ b_i' &= a_{2i} b_{2i-1} + b_{2i} \end{aligned}$$

2. Solve z_i recursively

3. For $1 \leq i \leq n$ set

$$\begin{aligned} y_i &= z_{i/2} && \text{if } i \text{ is even} \\ y_i &= a_i z_{(i-1)/2} + b_i && \text{if } i \text{ is odd } > 1 \\ y_i &= b_1 && \text{if } i = 1 \end{aligned}$$



First-Order Linear Recurrence

Input: Arrays $B = (b_1, b_2, \dots, b_n)$ and $A = (a_1 = 0, a_2, \dots, a_n)$, $n = 2^k$

Output: The y_i values such that $y_i = a_i y_{i-1} + b_i$

begin

if $n = 1$ **then** $y_1 := b_1$; **exit**

for $1 \leq i \leq n/2$ **pardo**

$a_i' := a_{2i} a_{2i-1}$

$b_i' := a_{2i} b_{2i-1} + b_{2i}$

Recursively solve the recurrence z_i defined by

$z_1 = b_1'$ and $z_i = a_i' z_{i-1} + b_i'$ for $2 \leq i \leq n/2$

for $1 \leq i \leq n$ **pardo**

if i is even **then** $y_i := z_{i/2}$

if $i > 1$ is odd **then** $y_i := a_i z_{(i-1)/2} + b_i$

if $i = 1$ **then** $y_1 := b_1$

end



Parallel Time and Work

- From the algorithm we observe
 - $T(n) = T(n/2) + O(1)$ therefore total parallel time $T(n) = O(\log n)$
 - $W(n) = W(n/2) + O(n)$ therefore total operations $W(n) = O(n)$



Polynomial Evaluation

- We wish to evaluate the polynomial

$$p(x) = b_1x^{n-1} + b_2x^{n-2} + b_3x^{n-3} + \dots + b_n$$

- Two steps:

1. Use prefix sum

- Compute the $x^{n-i} = [1, x, x^2, x^3, \dots, x^{n-1}]$ concurrently for all i , which takes $O(\log n)$ time and $O(n)$ work

2. Use a tree reduction to compute the sum

- Parallel sum $b_i x^{n-i}$ takes $O(\log n)$ parallel time and $O(n)$ work



Polynomial Evaluation (cont'd)

- We wish to evaluate the polynomial

$$p(x) = b_1x^{n-1} + b_2x^{n-2} + b_3x^{n-3} + \dots + b_n$$

- Horner's rule

$$p(x) = (((b_1 x + b_2) x + b_3) x + \dots + b_{n-1}) x + b_n$$

gives a first-order linear recurrence with $a_i = x$

- Takes $O(\log n)$ total parallel time with $O(n)$ total operations



Polynomial Multiplication

- Consider the polynomials

$$p(x) = a_0x^{n-1} + a_1x^{n-2} + a_2x^{n-3} + \dots + a_{n-1}$$
$$q(x) = b_0x^{m-1} + b_1x^{m-2} + b_2x^{m-3} + \dots + b_{m-1}$$

- We wish to compute the product

$$r(x) = p(x)q(x) = c_0x^{n+m-2} + c_1x^{n+m-3} + c_2x^{n+m-4} + \dots + c_{n+m-2}$$

where


$$c_k = \sum_{j=0..k} a_j b_{k-j}$$

(we take $a_j = 0$ for $j \geq n$ and $b_{k-j} = 0$ for $k-j \geq m$)



Polynomial Multiplication (cont'd)

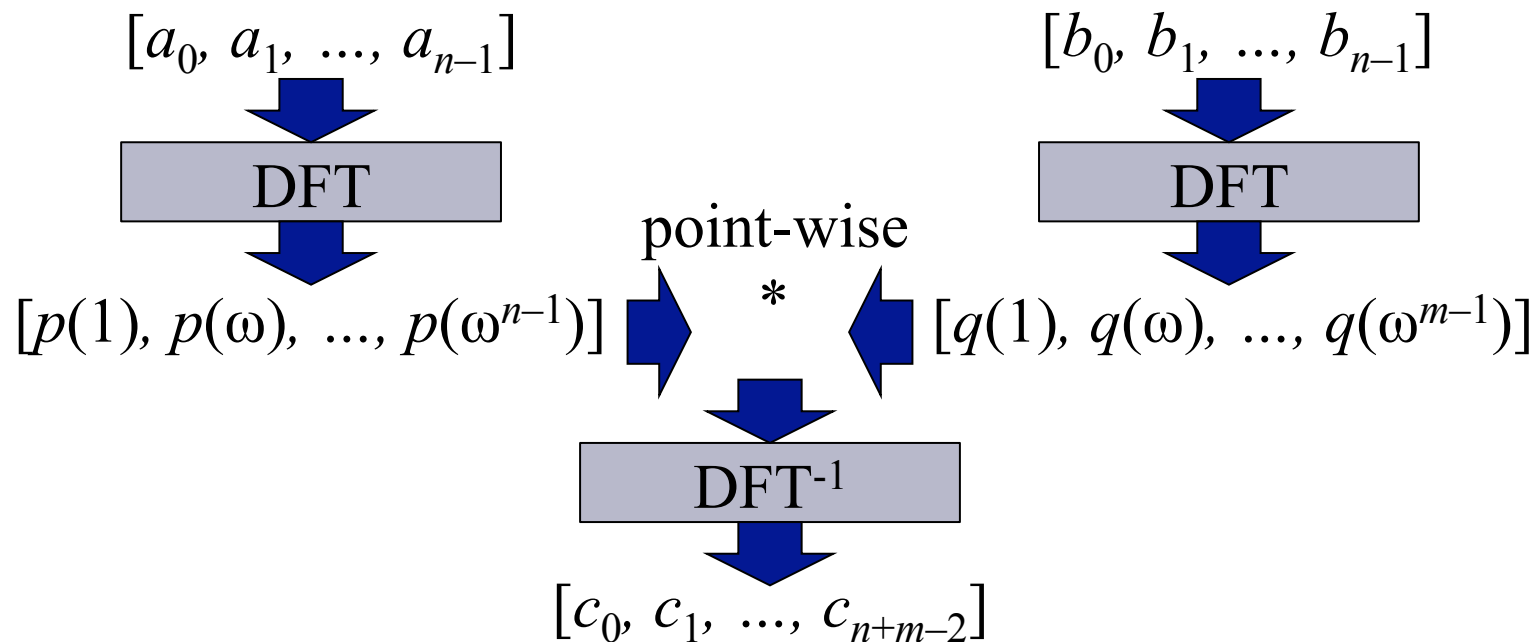
- We can compute all $c_k = \sum_{j=0..k} a_j b_{k-j}$ in $O(\log(n+m))$ parallel time
 - Takes $O((n+m-1)^2/2) = O(nm)$ operations





Polynomial Multiplication & FFT

- Convolution theorem: polynomial multiplication with FFT
 - $O(\log(n+m))$ parallel time, with a simple use of the FFT algorithm reduces the total number of operations to $O((n+m) \log(n+m))$
 - The FFT of the coefficients a_i of p and a_j of q gives the values of the product polynomial $r(\omega^j) = p(\omega^j)q(\omega^j)$ at the distinct roots of unity ω^j





Polynomial Multiplication & FFT

Input: Polynomial coeff. $\mathbf{a} = (a_0, a_1, \dots, a_{n-1})$ and $\mathbf{b} = (b_0, b_1, \dots, b_{m-1})$

Output: $\mathbf{c} = (c_0, c_1, \dots, c_{n+m-2})$ such that $c_k = \sum_{j=0..k} a_j b_{k-j}$

begin

1. Find integer $l = 2^s$ such that $n + m - 2 < l \leq 2(n + m - 2)$
2. Use FFT to compute $\mathbf{y} = \text{DFT}_l(\mathbf{a})$ and $\mathbf{z} = \text{DFT}_l(\mathbf{b})$
3. Compute $u_j = y_j z_j$ for all $j = 0, \dots, l-1$
4. Use FFT^{-1} to compute $\mathbf{c} = \text{DFT}_l^{-1}(\mathbf{u})$ giving $\mathbf{c} = (c_0, c_1, \dots, c_{l-1})$

end

Steps 2, 4 take $O(\log(n + m))$ parallel time and $O((n + m) \log(n + m))$ operations


Step 3 takes $O(1)$ parallel time and $O(n + m)$ total operations

Note: \mathbf{a} , \mathbf{b} , and \mathbf{c} vectors are implicitly padded with 0s, e.g. $a_i = 0$ for all $i \geq n$



Parallel FFT

- The FFT is easily parallelizable, since the fast sequential algorithm reduces the $O(n^2)$ problem into a $O(n \log n)$ problem using a divide-and-conquer strategy





Parallel FFT

Input: $\mathbf{x} = (x_0, x_1, \dots, x_{n-1})$, $n = 2^k$, $\omega = e^{i2\pi/n}$, where $i = \sqrt{-1}$
Output: $\mathbf{y} = \text{DFT}_n(\mathbf{x})$
begin
1. **if** $n = 2$ **then**
 $y_1 := x_1 + x_2$; $y_2 := x_1 - x_2$; **exit**
2. **for** $0 \leq j \leq n/2 - 1$ **pardo**
 $u_j := x_j + x_{n/2+j}$
 $v_j := \omega^j (x_j - x_{n/2+j})$
3. Recursively compute $\mathbf{z} := \text{DFT}_{n/2}(\mathbf{u})$ and $\mathbf{z}' := \text{DFT}_{n/2}(\mathbf{v})$
4. **for** $0 \leq j \leq n - 1$ **pardo**
 if j is even **then** $y_j := z_{j/2}$
 if j is odd **then** $y_j := z'_{(j-1)/2}$
end



Lagrange Interpolation

- Given a set of n points $\{(\alpha_j, \beta_j)\}_{j=0..n-1}$ determine the polynomial p of degree $n - 1$ such that for all $j = 0, \dots, n-1$

$$p(\alpha_j) = \beta_j$$

- Lagrange interpolation specifies p as follows

$$p(x) = \sum_{j=0}^{n-1} \beta_j \frac{\prod_{l=0, l \neq j}^{n-1} (x - \alpha_l)}{\prod_{l=0, l \neq j}^{n-1} (\alpha_j - \alpha_l)}$$



Lagrange Interpolation (cont'd)

- Divide-and-conquer strategy: rearrange terms
- Define

$$q_l = x - \alpha_l$$

and

$$Q(x) = \prod_{l=0}^{n-1} q_l = \prod_{l=0}^{n-1} (x - \alpha_l)$$

then the derivative of Q at point α_j is

$$Q'(\alpha_j) = \prod_{l=0, l \neq j}^{n-1} (\alpha_j - \alpha_l)$$

which can be evaluated $\gamma_j = Q'(\alpha_j)$, $c_j = \beta_j / \gamma_j$ giving

$$p(x) = \sum_{j=0}^{n-1} \beta_j \frac{Q(x) / (x - \alpha_j)}{Q'(\alpha_j)} = Q(x) \sum_{j=0}^{n-1} \frac{c_j}{x - \alpha_j}$$



Lagrange Interpolation (cont'd)

- A balanced tree can be used to compute the sum in

$$p(x) = Q(x) \sum_{j=0}^{n-1} \frac{c_j}{x - \alpha_j}$$

and use FFT-based polynomial multiplication

- There is another way: note that

$$\frac{p(x)}{Q(x)} = \frac{p_{k-1,0}(x)}{Q_{k-1,0}(x)} + \frac{p_{k-1,1}(x)}{Q_{k-1,1}(x)} = \frac{p_{k-1,0}(x)Q_{k-1,1}(x) + p_{k-1,1}(x)Q_{k-1,0}(x)}{Q(x)}$$

where

$$p_{k-1,0}(x) = Q_{k-1,0} \sum_{j=0}^{n/2-1} \frac{c_j}{x - \alpha_j}$$

$$p_{k-1,1}(x) = Q_{k-1,1} \sum_{j=n/2}^{n-1} \frac{c_j}{x - \alpha_j}$$

$$Q_{k-1,0}(x) = \prod_{j=0}^{n/2-1} q_l(x)$$

$$Q_{k-1,1}(x) = \prod_{j=n/2}^{n-1} q_l(x)$$



Lagrange Interpolation (cont'd)

Input: Set of pairs (α_j, β_j) for $j = 0, \dots, n-1$, $n = 2^k$

Output: The n coefficients of $p(x) = p_{k,0}(x)$ such that $p(\alpha_j) = \beta_j$

begin

1. **for** $0 \leq j \leq n-1$ **pardo**

$$Q_{0,j}(x) := x - \alpha_j$$

2. **for** $h = 1$ **to** $\log n$ **do**

for $0 \leq j \leq n/2^h - 1$ **pardo**

$$Q_{h,j}(x) := Q_{h-1,2j}(x) \times Q_{h-1,2j+1}(x)$$

3. Compute $Q'_{0,j}(x)$ and $\gamma_j := Q'_{0,j}(\alpha_j)$ for all $j = 0, \dots, n-1$

4. **for** $0 \leq j \leq n-1$ **pardo**

$$p_{0,j}(x) := \beta_j / \gamma_j$$

5. **for** $h = 1$ **to** $\log n$ **do**

for $0 \leq j \leq n/2^h - 1$ **pardo**

$$p_{h,j}(x) := p_{h-1,2j}(x) \times Q_{h-1,2j+1}(x) + p_{h-1,2j+1}(x) \times Q_{h-1,2j}(x)$$

end

$T(n)$ $W(n)$

$O(1)$ $O(n)$

$\left. \begin{array}{c} O(\log^2 n) \\ O(n \log^2 n) \end{array} \right\} O(\log n) \quad O(n^2)$


$O(1)$ $O(n)$

$\left. \begin{array}{c} O(\log^2 n) \\ O(n \log^2 n) \end{array} \right\}$



Convex Hull Problem Revisited


- The *planar convex hull* of a set of points $S = \{p_1, p_2, \dots, p_n\}$ of $p_i = (x, y)$ coordinates is the smallest convex polygon that encompasses all points S on the x - y plane





Convex Hull Problem Revisited

- The *upper convex hull* spans points $\{q_1, \dots, q_s\} \subseteq S$ from point q_1 with minimum x to q_s with maximum x
- The *convex hull* = *upper convex hull* + *lower convex hull*
- Problem:
 - Given points $S = \{p_1, \dots, p_n\}$ such that $x(p_1) < x(p_2) < \dots < x(p_n)$, construct the upper convex hull in parallel






Convex Hull Problem Revisited

- Points $S = \{p_1, \dots, p_n\}$ may have duplicate x-coordinate values
- Sort the points $x(p_1) \leq x(p_2) \leq \dots \leq x(p_n)$ in $O(\log n)$ parallel time and $O(n \log n)$ operations (pipelined merge sort)
- Then, if two or more points have the same x coordinate:
 - Keep the point with the largest y coordinate for the UCH
 - Keep the point with the smallest y coordinate for the LCH
- We can now assume that $x(p_1) < x(p_2) < \dots < x(p_n)$ to compute the UHS (and similarly the LHS)



Convex Hull Problem Revisited

- Parallel convex hull:
 1. Divide the x -sorted points S into sets S_1 and S_2 of equal size
 2. Compute upper convex hull recursively on S_1 and S_2
 3. Combine $UCH(S_1)$ and $UCH(S_2)$ by computing the upper common tangent to form $UCH(S)$





Convex Hull Problem Revisited

- Base case of recursion: two points, which are returned as $UCH(S)$
- *Revisit the common tangent computation:*
 - The line segment (a,b) can be computed sequentially in $O(\log n)$ time with $n = |UCH(S_1) + UCH(S_2)|$ using a binary search method
- *And replace with parallel computation:*
 - The line segment (a,b) can be computed in $O(1)$ parallel time
- Line segments can be implemented as linked list of points, thus $UCH(S_1)$ and $UCH(S_2)$ can be connected using one pointer change of a to point to b in $O(1)$ time




Intermezzo: Parallel Search

- Let $X = (x_1, x_2, \dots, x_n)$ be n distinct elements from a set S such that $x_1 < x_2 < \dots < x_n$
- Given $y \in S$, find the index i for which $x_i \leq y < x_{i+1}$ where we added $x_0 = -\infty$ and $x_{n+1} = +\infty$
- Parallel search with p processors:
 - Split X in p segments of (almost) equal length
 - Each processor verifies if y is in its segment
 - If so, restrict search to the segment containing y and repeat



Intermezzo: Parallel Search

2	4	6	8	10	12	14	16	18	20	22	24	26	28	30
---	---	---	---	----	----	----	----	----	----	----	----	----	----	----




$$\text{Parallel time} = O\left(\frac{\log(n+1)}{\log(p+1)}\right)$$

p_2 found $i=13$



Convex Hull Problem Revisited: Using Parallel Search


- Let $\text{UCH}(S_1) = (r_1, \dots, r_s)$ and $\text{UCH}(S_2) = (q_1, \dots, q_t)$
- We need to determine points $a = r_i$ and $b = q_{j(i)}$ such that all points in S are below the line through points a and b





Convex Hull Problem Revisited: Using Parallel Search


- Given a point $r_i \in \text{UCH}(S_1)$ then for any $q_k \in \text{UCH}(S_2)$ we can determine in $O(1)$ sequential time if $q_k = q_{j(i)}$, or $q_{j(i)}$ is to the left of q_k or $q_{j(i)}$ is to the right of q_k
- Thus, using parallel search, we can determine for point r_i the tangent $(r_i, q_{j(i)})$ in $O(\log t / \log p)$ parallel time using p processors, where $t = |\text{UCH}(S_2)|$





Convex Hull Problem Revisited: Using Parallel Search

- Given a point $r_i \in \text{UCH}(S_1)$ and $q_{j(i)} \in \text{UCH}(S_2)$, then we can determine in $O(1)$ sequential time if $r_i = a$, or a is to the left of r_i or a is to the right of r_i
- Thus, using parallel search, we can determine the tangent (a, b) in $O(\log(st) / \log p)$ parallel time using p processors, where $s = |\text{UCH}(S_1)|$ and $t = |\text{UCH}(S_2)|$





Convex Hull Problem Revisited: Using Parallel Search

- Take $p = \sqrt{s}\sqrt{t}$ then
 $O(\log(st) / \log(\sqrt{s}\sqrt{t})) = O(\log(st) / \frac{1}{2}\log(st)) = O(1)$
parallel time and $O(\sqrt{s}\sqrt{t}) = O(n)$ operations
- 1. Choose \sqrt{s} points from $UCH(S_1)$ thereby dividing the set $UCH(S_1)$ into (almost) equal blocks of size \sqrt{s} each
- 2. Find the $q_{j(k\sqrt{s})}$ for each $r_{k\sqrt{s}}$, $k = 1, \dots, \sqrt{s}$, using $p = \sqrt{s}\sqrt{t}$ processors in $O(1)$ parallel time
- 3. Deduce the block $B_k = (r_{k\sqrt{s}+1}, \dots, r_{(k+1)\sqrt{s}-1})$ that contains a
- 4. For each r_i in block B_k , determine $q_{j(i)}$ and search $a = r_i$ using $p = \sqrt{s}\sqrt{t}$ processors in $O(1)$ parallel time
- 5. Set $b = q_{j(i)}$



Convex Hull Problem Revisited: Putting it Together

- Preprocess the points by sorting in $O(\log n)$ parallel time (pipelined merge sort), such that $x(p_1) \leq x(p_2) \leq \dots \leq x(p_n)$
- Remove duplicates $x(p_i) = x(p_j)$ (for UCH and LCH)
- Divide-and-conquer:
 1. Split S into S_1 S_2 and recursively compute the UCH of S_1 and S_2
 2. Combine $UCH(S_1)$ and $UCH(S_2)$ by computing the upper common tangent in $O(1)$ time to form $UCH(S)$
- Repeat to compute the LCH
- Parallel time (assuming $p = O(n)$ processors)

$$T(n) = T(n/2) + O(1)$$

gives

$$T(n) = O(\log n)$$



Further Reading

- An Introduction to Parallel Algorithms, by J. JaJa, 1992